

Future Suffering and the Non-Identity Problem

Theron Pummer (University of St Andrews)

Global Priorities Institute | July 2024

GPI Working Paper No. 17-2024

Please cite this working paper as: Pummer, T. *Global Priorities Institute Working Paper Series, No. 17-2024*. Available at <https://globalprioritiesinstitute.org/future-suffering-and-the-non-identity-problem-theron-pummer>



Future Suffering and the Non-Identity Problem¹

Theron Pummer

tgp4@st-andrews.ac.uk

Abstract: I present and explore a new version of the *Person-Affecting View*, according to which reasons to do an act depend wholly on what would be said for or against this act from the points of view of particular individuals. According to my view, (i) there is a morally *requiring* reason not to bring about lives insofar as they contain suffering (negative welfare), (ii) there is *no* morally requiring reason to bring about lives insofar as they contain happiness (positive welfare), but (iii) there is a *permitting* reason to bring about lives insofar as they contain happiness. I show how my view solves the non-identity problem, while retaining the procreation asymmetry and avoiding implausible forms of antinatalism. We can be morally required to ensure that the quality of life of future people is higher rather than lower when this involves bringing about (worth living) lives that would contain less suffering rather than bringing about different (worth living) lives that would contain more suffering.

1. Introduction

Let's temporarily set aside the effects of creating new people on those who will exist regardless of what we choose. Are we ever morally required to create new people, *just because* these new people would live happy lives? Does the fact that there could be trillions

¹ For helpful written comments, I am grateful to Jonas H. Aaron, Krister Bykvist, Matthew Clark, Owen Clifton, Roger Crisp, Tomi Francis, Elizabeth Harman, Jonas Harney, Joe Horton, Joel Joseph, Rahul Kumar, Jeff McMahan, Andreas Mogensen, Melinda Roberts, Thomas Sinclair, Luca Stroppa, Christian Tarsney, Teruji Thomas, Elliott Thornley, and Travis Timmerman. I am also grateful to audiences at the University of St Andrews, Seton Hall University, the Institute for Futures Studies (Stockholm), Lund University, and the University of Oxford. An earlier version of this paper was given as the Derek Parfit Memorial Lecture at the University of Oxford on 12 June 2024.

upon trillions of happy people in the far future itself provide us with overwhelming moral pressure to avoid extinction?²

According to *Total Utilitarianism*, you are always morally required to maximize total welfare. This view has highly counterintuitive implications. If doing so maximizes total welfare, Total Utilitarianism says you'd be morally required to allow enormous costs to befall each presently existing person to ensure that the far future contains very many happy people. Worse still, it says you'd be morally required to allow these huge costs to befall existing people to ensure that there are 10^{40} happy people in the far future, even if there otherwise would have been "only" 10^{20} happy people.

I take a very different view. Not only do I reject Total Utilitarianism, but I reject *any* view that implies you are morally required to create new people, just because these new people would live happy lives. I do not deny that there is moral pressure to avoid extinction. I only deny that the happiness of possible future people is a source of moral pressure to avoid extinction.

I will defend a version of the *Person-Affecting View*, according to which whether an act is morally permissible depends wholly on the way it would affect particular people. The guiding question behind this view is this: *Would there be something to be said against your act, from the self-interested point of view of any particular person?*

If you fail to rescue a drowning stranger, there would be something to be said against your act, from this very stranger's point of view. The same does not hold in the case of failing to create happy people. To isolate what matters, consider an artificially controlled example.

² Many have answered in the affirmative, including Glover (1977, 69), Parfit (1984) (2011), Bostrom (2013), and Beckstead (2019).

Happy Eve: You can press a button that will, at no cost or benefit to anyone who exists independently, cause Eve to pop into existence. Her life would contain 60 years of happiness and no suffering.

There's nothing to be said against not creating Eve, from Eve's point of view. This is because, if you fail to create Eve, she will never exist. There can be no "ghostly" complaints against not creating happy people.³ So, according to the Person-Affecting View, the fact that the lives of possible future people would contain happiness does not generate a moral requirement to create them. This does not mean the Person-Affecting View simply *ignores* the welfare of merely possible people. That would be a mistake. To see this, consider another case.

Miserable Zeke: You can press a button that will, at no cost or benefit to anyone who exists independently, cause Zeke to pop into existence. His life would contain 60 years of suffering and no happiness.

Clearly, it is wrong to create Zeke. There *is* something to be said against creating Zeke, from Zeke's point of view. After all, if you create him, he will have a miserable existence. So, according to the Person-Affecting View, the fact that the lives of possible future people would contain suffering *does* generate a moral requirement *not* to create them.

The Person-Affecting View supports what is known as the

³ This expression is from Podgorski (2023, 348).

Procreation Asymmetry: There's no moral requirement to create happy people (in virtue of their being happy), but there *is* a moral requirement *not* to create miserable people (in virtue of their being miserable).⁴

Many find each “half” of the Procreation Asymmetry to be plausible, and the Person-Affecting View appears to provide a decent explanation.⁵

However, Derek Parfit famously presented a major problem for the Person-Affecting View. It is known as the *non-identity problem*. To bring out the non-identity problem, consider an example concerning climate change.⁶

Carbon Policy: We can carry on with business as usual or adopt a widespread policy of reducing carbon emissions. While this carbon policy would be slightly worse for each of the 8 billion of us alive now, it would result in a *much* higher level of welfare for the trillions of people who would exist two centuries from now and beyond.

However, adopting the carbon policy would also make it the case that the trillions who would exist in two centuries are *numerically different people than* the trillions who would have existed if we had instead carried on with business as usual.

Compared to the status quo, the adoption of the carbon policy would change the timings and manners of conceptions, changing who would exist in the future. These

⁴ For discussion, see Narveson (1973), McMahan (1981), Roberts (2011), Chappell (2017), Frick (2020), and Spencer (2021).

⁵ One might object that the Person-Affecting View is itself appealing to an asymmetry between personal complaints and personal gratitude. Roughly: an act is wrong only if it would give rise to a complaint, so that failing to act in a way that would give rise to gratitude isn't wrong in the absence of any complaint. But many find this complaints-gratitude asymmetry built into the Person-Affecting View quite a compelling way of thinking about moral wrongness. On complaint-based views, see Parfit (1984, 526), Scanlon (1998), McDermott (2019), Horton (2021a), and Podgorski (2023). If we buy into such a view, it looks like we can provide a decent explanation of the Procreation Asymmetry.

⁶ Based on Parfit's *Depletion* (1984, 361).

effects would snowball. No one who would exist in two centuries, if we carry on with business as usual, would exist in two centuries, if we adopt the carbon policy.

It seems clear that we are morally required to adopt the carbon policy. And it seems clear that we are morally required to do so *because* this would result in a much higher quality of life for future people. The Person-Affecting View struggles to capture this intuition. *Would* there be anything to be said against carrying on with business as usual, from the point of view of any particular person? The description of the case stipulates that carrying on with business as usual would be better for each of us alive now, so it looks like there'd be nothing to be said against doing so from the point of view of any one of us. What about the people who would exist in two centuries? Well, if their lives were miserable, so that they were overall worth *not* living, then there'd be something to be said against carrying on with business as usual from the points of view of these future people. *But suppose that this is not the case*: suppose that the future people who'd exist if we carried on with business as usual would have overall happy lives. *Still*, it seems we'd be morally required to adopt the carbon policy, on the grounds that this would result in a much higher quality of life for future people.

It appears that the Person-Affecting View cannot capture this intuition. It looks like there'd be nothing to be said against business as usual from the point of view of any one of the particular people who would exist if we carried on with business as usual. After all, if we had instead adopted the carbon policy, these people would never exist. This is one instance of the non-identity problem.

In this paper, I will present and explore a new version of the Person-Affecting View that can solve the non-identity problem. My view can explain, for example, why we are morally required to adopt the carbon policy rather than carry on with business as usual. What

is more, my view retains the Procreation Asymmetry and avoids implausible forms of antinatalism.

2. Setup

A bit of setup is needed before I can properly state my version of the Person-Affecting View. First, some terminological points. I will take “happiness” to refer to positive welfare, well-being, or whatever contributes toward making a life worth living, be it pleasure, desire satisfaction, achievement, friendship, or some combination of these things (welfare components may exhibit various contextual interactions). For simplicity, I’ll illustrate the welfare levels of lives in terms of years of happiness. Similarly, I will take “suffering” to refer to negative welfare, ill-being, or whatever contributes toward making a life worth not living. I take a “happy life” to be one that is overall worth living, containing an overall positive balance of welfare components, and a “miserable life” to be one that is overall worth *not* living.

Much of population ethics is concerned with *axiology*. That is, it is concerned with ranking possible outcomes according to their *intrinsic value*. Compare two possible outcomes:

- (1) One billion happy people, each at welfare level 100.
- (2) Two billion *different* happy people, each at welfare level 200.

Is (2) *better* than (1)? Many will say that it is. After all, the total and average welfare of (2) is greater than that of (1). However, not everyone agrees. Some have defended the Person-Affecting View as an *axiological* view.⁷ According to a *Strong Axiological Person-Affecting View*, one possible outcome is *better* than only if it is *better for* at least one particular person.

⁷ For discussion, see (for example) Temkin (2012) and Ross (2015).

Since no one who exists in (2) also exists in (1), (2) is not better for any particular person than (1). So, according to the Strong Axiological Person-Affecting View, (2) is not better than (1).

I will not be defending an Axiological Person-Affecting View. Instead, I will be defending the Person-Affecting View as a *deontic* view. It is concerned with the *deontic status* of acts, that is, whether these acts are morally wrong, permissible, or required. I take it that an act is permissible if and only if it is not wrong, and that an act is required if and only if it is wrong not to do it. The Person-Affecting View I defend says that the moral permissibility of an act depends wholly on whether there would be something to be said against your act, from the point of view of any particular person. It does *not* imply that (2) is *not better* than (1). It simply makes no comment on the intrinsic value of outcomes.

What's more, I will be assuming *thoroughgoing non-consequentialism*, according to which there is no moral pressure to promote intrinsic value *per se*.⁸ The fact that an act would make the world a better place cannot itself generate a morally requiring reason to perform that act. So, while I won't be commenting on the intrinsic value of outcomes, *even if* beaming happy people into existence on some distant uninhabited planet *did* make the world a better place, it wouldn't follow that you'd be under any moral pressure to make the world a better place by creating these happy people.

It may at first seem shocking to suggest that the fact that an act would make the world a better place cannot itself ground a moral requirement to perform that act. Surely, you're morally required to rescue a drowning child if the only cost were muddied shoes—and presumably rescuing the child would make the world a better place! Well, I claim you *are* required to save the child, but not *because* doing so would make the world a better place.

⁸ I borrow this term from Sinclair (2018). Also see Scanlon (1998, 222).

Your requirement to save the child is instead grounded in the fact that there would be something to be said, from this child's point of view, against letting them die.

My primary focus is the permissibility of acts. I will make use of a framework of reasons that *determine* the permissibility of acts. The first key feature of this framework is the distinction between morally requiring reasons and morally permitting reasons.⁹

A *morally requiring reason* to perform an act contributes toward making this act morally required. Absent countervailing considerations, it is wrong not to do this act. Morally requiring reasons can conflict. For example, there is a morally requiring reason not to shove me onto the ground. There is a stronger morally requiring reason to save a stranger's life. If shoving me onto the ground is the only way to save an innocent stranger's life, then there is an *overall* morally requiring reason to shove me rather than not.

By contrast, a *morally permitting reason* to do an act contributes toward making this act morally permissible without contributing toward it morally required. It has morally permitting strength, which contributes toward defeating the wrong-makingness supplied by countervailing morally requiring reasons. A morally permitting reason to do an act is still (I claim) a *reason* in that it counts in favor of doing the act. It has both a favoring role and a permitting role.

In my book *The Rules of Rescue* I use the framework of requiring reasons and permitting reasons to explain the permissibility of failing to help strangers at great cost to yourself. Compare two cases.

⁹ The distinction between requiring reasons and permitting reasons comes from Gert (2003). Gert distinguishes between a reason's rationally requiring strength and its rationally justifying strength. My interest here is in the distinction between *morally* requiring strength and *morally* justifying strength. For an explication of the latter distinction, see Portmore (2011, 121).

Costless Rescue: You can at *no* cost to yourself prevent a stranger from being crushed to death by a boulder. By effortlessly pressing a button, you can raise a wall, stopping the boulder before it reaches the stranger.

Costly Rescue: You can at *great* cost to yourself prevent a stranger from being crushed to death by a boulder. You can save the stranger only if you place your legs in the boulder's path, which would result in the loss of your legs.

I take it that, in both cases, there is a morally requiring reason to save the stranger's life. In *Costless Rescue*, there are no other considerations in play. There's an overall requiring reason to save the stranger's life, and no permitting reason not to. You are therefore morally required to save the stranger's life. In *Costly Rescue*, there's more at stake than the stranger's life. Your legs are also at stake. Still, I take it that there is an *overall* morally requiring reason to save the stranger rather than refrain. Crucially, however, there is a morally *permitting* reason not to sacrifice your legs, and it is plausible that it is *sufficiently strong* even in the face of the strong overall requiring reason to save the stranger's life. You are therefore permitted not to save the stranger, but you are also not required *not* to save the stranger. Each act—saving and not saving—is morally permissible.

If the cost of saving the stranger were quite small, say, the muddying of your shoes, then the morally permitting reason to avoid incurring this cost would be insufficiently strong next to the overall requiring reason to save a stranger's life. The weak permitting reason not to muddy your shoes would still *contribute* toward defeating the wrong-makingness supplied by the overall requiring reason to save the stranger, but it would not actually succeed in defeating it. When the cost of helping is relevantly small, it's wrong not to help.

The second key feature of my framework of reasons is a principle that links reasons with wrongness.

Reasons-Wrongness Link: It is wrong to do an act *A* if and only if there is an overall morally requiring reason to do an alternative *B* instead, and insufficient permitting strength supporting this act *A* over the alternative *B*.

The Reasons-Wrongness Link is a pairwise-tournament choice rule, in the following sense. An alternative *A* is “defeated” by an alternative *B* if there’s an overall morally requiring reason to do *B* instead of *A*, and no sufficient permitting strength supporting *A* over *B*. If an alternative is defeated by *any* other alternative, then it’s wrong (it loses the tournament, so to speak). This is true even if this defeated alternative is undefeated by many other alternatives. To see how this works, consider another rescue case.

Costly No-Conflict: Strangers Albert and Betty face a deadly threat. You have three alternatives. You can do nothing, in which case both strangers will die. You can save just Albert’s life at great cost to yourself. You can save both Albert *and* Betty at the same great cost to yourself.

I’ve elsewhere argued that, while the alternative of saving Albert only isn’t defeated by the alternative of doing nothing, it *is* defeated by the alternative of saving both Albert and Betty.¹⁰ There’s an overall requiring reason to save both rather than save Albert only, and no

¹⁰ Pummer (2023, chapter 3). For relevant discussion, see Horton (2017), Rulli (2020), Muñoz (2021), and Tucker (forthcoming).

sufficiently strong permitting reason to save Albert only rather than save both. By the Reasons-Wrongness Link, it is wrong to save Albert only.

This is all just meant to get you familiar with the framework of requiring reasons, permitting reasons, and the Reasons-Wrongness Link. I will now put this framework to use in population ethics.

3. A New Version of the Person-Affecting View

I've already given a loose articulation of the Person-Affecting View. I will now begin to offer a more precise interpretation of the Person-Affecting View. Here's the general idea:

Person-Affecting View: The reasons to do an act depend wholly on what would be said for or against this act from the self-interested points of view of particular individuals.

As I interpret it, this view recognizes *requiring* reasons based on complaints: there is a requiring reason to do act *A* if and only if there would be something to be said *against* not-*A* from some particular individual's point of view. It also recognizes *permitting* reasons based on gratitude: there is at least presumptively a permitting reason to do an act *A* if there would be something to be said in *favor* of *A* from some particular individual's point of view. I will say more about both sorts of reasons below. In section 7, I will go on to argue that my view should include *silencing* considerations, in addition to requiring and permitting considerations.

As I understand the Person-Affecting View, it is *deontic* (ultimately concerned with permissibility and wrongness) and *narrow* with respect to wrong-makingness (the basis of an act's wrongness is how *particular* individuals would be affected by it, and not with how

people in general would be affected, or with the promotion of intrinsic value).¹¹ To better understand the Person-Affecting View, it is useful to consider some examples.

Happy Amy: You can press a button that will, at no cost or benefit to anyone who exists independently, cause Amy to pop into existence. Her life would contain 60 years of happiness *and 6 months of suffering*. Overall, her life is worth living. (In the not-to-scale figures below, green is for happiness and red is for suffering.)



What does the Person-Affecting View say about this case? It implies there is no requiring reason *to* create Amy. As we've already seen, there would be nothing to be said against not creating Amy, from Amy's point of view. The more interesting question is whether the Person-Affecting View implies that there is a requiring reason *not* to create Amy.

Whether the Person-Affecting View implies there is a requiring reason not to create Amy depends on whether the phrase "*something to be said against*" should be interpreted in an *overall* way or in a *pro tanto* way. Since her life would be overall worth living, Amy would not have an overall complaint against being created. Nonetheless, the Person-Affecting View should be concerned with the fact that Amy would have a *pro tanto* complaint: the

¹¹ What could be said against your act, from the self-interested point of view of a particular individual, can include more than what would be *prudentially worse* for this individual. For example, the fact that an act leaves me worse off *than others*, or the fact that it *causes* me harm rather than merely allows me to suffer harm, could give me a stronger individual complaint against this act. Relatedly, see Scanlon (1998, 219-20) on the bases of personal complaints.

suffering in Amy's life is something to be said against creating Amy, from Amy's point of view. So, there is a requiring reason not to create Amy.

Suppose Zara, by contrast, would have 6 months of suffering and 1 month of happiness. Her life would be worth *not* living. Zara would have an overall complaint against being created. Does the 6 months of suffering in Zara's life provide a stronger requiring reason not to create Zara, than the 6 months of suffering in Amy's life provides not to create Amy? Or do they provide equally strong requiring reasons? For now, I shall leave this an open question, though there is something intuitive about the idea that Zara's 6 months of suffering provides an even stronger requiring reason against creation.

Putting together what I've said so far, my version of the Person-Affecting View implies the following two claims about requiring reasons:

- (i) There is a requiring reason not to bring about lives insofar as they contain suffering (the strength of this requiring reason is an increasing function of the magnitude of this suffering).
- (ii) There is no requiring reason to bring about lives insofar as they contain happiness.

According to claims (i) and (ii), there is a requiring reason not to create Amy and no requiring reason to create her. More generally, these claims imply that, if X's life would contain any suffering, there is a requiring reason not to create X, but no requiring reason *to* create X, no matter how much happiness X's life would contain.

It would appear that claims (i) and (ii) together imply:

Same-Life Antinatalism: When your only two alternatives are to create X or create no one, it is wrong to create X just in case X's life contains *any* suffering—even if X's life is overall *well* worth living.

It is one thing to claim that it is wrong to have children because of the bad effects this would have on those who exist already, it is quite another to claim that it is wrong to create an overwhelmingly happy person just because their life would contain a tiny bit of suffering. Indeed, I take it that many would hold that it is morally permissible to create Amy in *Happy Amy*, even though her life contains *6 months* of suffering. Although it has some notable defenders, I find Same-Life Antinatalism deeply implausible.¹²

In addition to claims (i) and (ii), we should accept,

- (iii) There is a *permitting* reason to bring about lives insofar as they contain happiness. That is, there is a permitting reason to create someone *rather than* create someone else or create no one at all, insofar as the former alternative brings about *more* happiness than the latter alternative (the strength of this permitting reason is an increasing function of the magnitude of the difference in happiness).

Claim (iii) is implied the Person-Affecting View, suitably interpreted. After all, if you create a happy person, there would be something to be said *in favor of* your act from this particular person's point of view—and there'd more to be said in favor of your act from their point of view, the more happiness their life would contain. However, this element of the Person-Affecting View is not "narrow" in that there is only *presumptively* a permitting reason to do

¹² One such defender is Benatar (2006).

an act *A* when there would be something to be said in favor of *A* from some particular individual's point of view. I take it that there is *no* permitting reason to create happy Amy *rather than* create happy Bobby, if Amy's life would contain *no more* happiness than Bobby's. As I explain below, this contrastive aspect of (iii) is one of the keys to solving the non-identity problem.¹³

Permitting reasons to create people insofar as their lives contain happiness have a *favoring* role and a *permitting* role. In this paper, I will largely focus on their permitting role, but let me first briefly say something about their favoring role. Creating happy people would give rise to gratitude from those created. This provides a reason in favor of creating happy people. It is a merely favoring reason, or "enticing" reason, in that it does not contribute toward making the favored act morally required—after all, not creating happy people would not give rise to any complaints.¹⁴ Not only would it not be *wrong* not to respond to this reason to create happy people, but I also claim it would not be in any way mean, unkind, or irrational. Just as there can be no "ghostly" complaints against not creating happy people, there can be no "ghostly" negative residue of not responding to merely favouring reasons to create happy people. Things may be different in the case of merely favoring reasons to do nice things for existing people; perhaps it is somewhat unkind or irrational not to do these nice things for them, when you can do so entirely costlessly. While I find these thoughts about the favoring role of permitting reasons to create happy people to be plausible, they are not as crucial for the main claims I wish to defend here.

What's most crucial for my purposes here is the *permitting* role of these permitting reasons: the fact that an act creates people whose lives contain happiness can make the act permissible even when it is opposed by significant requiring reasons. When we combine

¹³ For a general account of contrastivism about reasons, see Snedegar (2017).

¹⁴ On non-requiring but (still) favoring reasons, see Dancy (2021), Little and Macnamara (2017), and Horgan and Timmons (2010).

claims (i) and (ii) with claim (iii)—the claim invoking permitting reasons—we avoid Same-Life Antinatalism. For example, while the 6 months of suffering in Amy’s life provides a requiring reason not to create her, it is plausible that the 60 years of happiness in her life provides a sufficiently strong permitting reason *to* create her. That is, the happiness provides a strong enough permitting reason to make it permissible to create her, even though there is a significant requiring reason not to create her. (If you disagree that 60 years of happiness is enough to make it morally permissible to justify 6 months of suffering, you can change the example by increasing the happiness and/or decreasing the suffering in Amy’s life; however, it is important to note that the stricter you are about these proportions, the closer you will get to Same-Life Antinatalism!)

In real life, there is uncertainty. Consider the following variant:


Expectably Happy Amy: You can press a button that will, at no cost or benefit to anyone who exists independently, cause Amy to pop into existence. There is a 99% chance her life will contain 60 years of happiness and no suffering, and a 1% risk her life will contain 6 months of suffering and no happiness.

One might reason as follows. Creating Amy carries a 99% chance of creating a happy person, which there is no requiring reason to do (but a permitting reason to do) and creating her carries a 1% risk of creating a miserable person, which there is requiring reason not to do. Therefore, creating Amy carries a 99% chance of doing something permissible but not required and a 1% risk of doing something wrong. So, deciding under uncertainty, you shouldn’t take this risk, you shouldn’t create Amy. We should reject this reasoning (which would imply a risk-based variant of Same-Life Antinatalism). Here is how I propose we approach *Expectably Happy Amy* instead: while the *small risk* of suffering itself provides a

relatively weak requiring reason not to create Amy, the *big chance* of happiness itself provides a strong permitting reason to create Amy, so that it is permissible to create her.¹⁵ I will from here onwards work with cases that involve no uncertainty.

Consider another kind of case.

Happy & Miserable: You can press a button that will cause *two* people to pop into existence, Albert and Zara. While Albert’s life would contain 60 years of happiness and no suffering, Zara’s life would contain 6 months of suffering and no happiness.



Many of us will find it wrong to press this button, even if we think it is permissible to create Amy in *Happy Amy*. Some might hold that, while Amy’s happiness justifies Amy’s suffering, Albert’s happiness *does nothing to justify* Zara’s suffering. Some might wish to impose such a same-life restriction on claim (iii).

However, a same-life restricted (iii) together with claims (i) and (ii) will imply,

Different-Lives Antinatalism: When your only two alternatives are to create the A-people & Z, or create no one, it is wrong to create the A-people & Z, if Z’s life is

¹⁵ I here follow Thomas (2023, sections 2.2 and 4.4). I later note in the main text that Thomas’s view crucially differs from mine in that, on his view, reasons are generated by (expected) overall welfare levels, rather than (expected) welfare components like happiness and suffering.

overall *barely* miserable—even if there are many A-people whose lives are overall *well* worth living.

This form of antinatalism seems implausibly extreme. Suppose you could, with the push of a button, cause billions of very happy A-people to pop into existence on some distant uninhabited planet, with the necessary side effect of bringing into existence one *only marginally* miserable Z-person. It seems to me permissible to push the button.

Different-Lives Antinatalism is not *as* implausible as Same-Life Antinatalism. The separateness of persons matters: if creating Amy with some suffering in her life is a necessary side effect of creating some amount of happiness, this fact provides a permitting reason to create her. But the permitting reason to create Amy seems *stronger* when this happiness is realized in Amy's life than when it's realized in the lives of others (and perhaps suffering generates stronger requiring reasons against creation when it's realized in lives worth not living than when it's realized in lives worth living). Still, we should reject an absolute same-life restriction on (iii). We should reject Different-Lives Antinatalism.

Now suppose Zeke's life would be filled with 60 years of suffering and no happiness. Could the fact that creating him is a necessary side effect of creating billions of *very happy* people provide a sufficiently strong permitting reason to create him? Though more controversial, I suspect the answer is Yes. This seems like a "one versus many" case. If you create miserable Zeke and the billions of very happy people, there would much to be said in favor of your act, taking account of the points of view of many happy individuals—I suspect this would be enough to justify what there is to be said against your act, from Zeke's point of view. It seems that the points of view of separate people in this way at least partially aggregate. However, it is not *intuitively* plausible that they fully aggregate: intuitively, it is

not permissible to create miserable Zeke as a necessary side effect of creating *marginally happy* people, no matter how numerous.¹⁶

To recap, the Person-Affecting View I have been developing includes three key claims:

- (i) There is a requiring reason not to bring about lives insofar as they contain suffering.
- (ii) There is no requiring reason to bring about lives insofar as they contain happiness.
- (iii) There is a permitting reason to bring about lives insofar as they contain happiness.

These three claims are motivated by foundational and unified person-affecting ideas, and they deliver plausible implications in a range of cases. For the sake of simplicity, in this paper I focus on cases of creating lives (“doing”) rather than cases of allowing lives to be created (“allowing”), but my view is intended to cover both sorts of cases.¹⁷

I’d like to briefly flag that my Person-Affecting View is importantly different from potentially seemingly related views in existing literature.

While Teruji Thomas’s view (2023) implies there are requiring reasons not to create miserable lives, and permitting reasons to create happy lives, his view doesn’t include claim (i). On his view, reasons are generated by overall welfare levels, rather than welfare components (happiness and suffering). For example, Thomas’s view does not recognize a requiring reason not to create Amy in *Happy Amy*. As we will see in section 5, claim (i) is needed to solve the non-identity problem.¹⁸

¹⁶ For literature on full versus partial aggregation, see Kamm (2007), Voorhoeve (2014), Horton (2021b), and Mann (2022).

¹⁷ My full view therefore includes the following claims:

(i*) There is a requiring reason not to allow lives to come into existence insofar as they contain suffering.

(ii*) There is no requiring reason to allow lives to come into existence insofar as they contain happiness.

(iii*) There is a permitting reason to allow lives to come into existence insofar as they contain happiness.

My view can be combined with the *Doctrine of Doing and Allowing*, so that for example the requiring reason implied by (i) is stronger than the requiring reason implied by (i*), other things being equal. For a defense of the Doctrine, see Woollard (2015).


¹⁸ Algander (2012) and Mogensen (2019) also draw the requiring/permitting distinction in the context of the Procreation Asymmetry, but, like Thomas, they do not include the crucial claim (i). Aaron (2024) draws the requiring/permitting distinction and does include claim (i), but in doing so he essentially appeals to a

Like my view, Elizabeth Harman's (2004) (2009) view implies there is a requiring reason not to create Amy in cases like *Happy Amy*. However, her view does not draw a distinction between requiring reasons and permitting reasons.¹⁹ Without such a distinction between different kinds of reasons, there is a danger we will recognize only requiring reasons (albeit of different strengths). If we have only requiring reasons to work with, we avoid antinatalism only if we claim that there are requiring reasons to bring about lives insofar as they contain happiness, contrary to the Procreation Asymmetry and the Person-Affecting View as I understand it.

4. Same-Life Improvement

While claims (i)-(iii) will do most of the relevant work in this paper, they are not sufficient. To see this, consider the following,

Same-Life Improvement: You have three alternatives. Create no one, create Amy with 40 years of happiness and no suffering, create Amy with 60 years of happiness and no suffering.



controversial "Shortfall View" (I raise doubts about such views in section 6). I show how we can adequately handle a variety of non-identity cases without appealing to this sort of view.

¹⁹ Nor do other harm-based views—such as those defended by Shiffrin (1999) and Gardner (2015)—draw the requiring/permitting distinction.

It seems wrong to create Amy with 40 years of happiness, but claims (i)-(iii) imply that it is permissible to create Amy with either amount of happiness. In addition to claims (i)-(iii), we should accept,

- (iv) There is a requiring reason to create a happy person with more happiness *rather than* create this *same* person with less happiness (the strength of this requiring reason is an increasing function of the magnitude of the difference in happiness).

Claim (iv) is not an ad hoc addition. It follows from the Person-Affecting View, according to which there is a requiring reason to do act *A* if and only if there would be something to be said against *not-A* from some particular individual's point of view. From Amy's point of view, there *would* be something to be said against creating her with 40 years of happiness *rather than* creating her with 60 years of happiness: you've made her worse off than you could have done. Given that there are no other requiring reasons in play, there is an overall requiring reason to create Amy with 60 rather than create her with 40. Moreover, there is no permitting reason to create Amy with 40 years of happiness rather than create her with 60 years of happiness. So, given the Reasons-Wrongness Link, it is wrong to create Amy with 40 years of happiness.

Provided that it includes claim (iv), the Person-Affecting View has the plausible implication that it is wrong to create Amy with 40 years of happiness in *Same-Life*

Improvement.²⁰ However, the inclusion of (iv) may lead to difficulties in the following sort of case.²¹

Amy, Bobby, and Carley: You have three alternatives. In the table below, numbers represent welfare levels and Ω represents nonexistence (i.e., the person in question isn't created in the alternative in question).

| Alternative | Amy | Bobby | Carley |
|-------------|----------|----------|----------|
| (A) | 100 | 50 | Ω |
| (B) | Ω | 100 | 50 |
| (C) | 50 | Ω | 100 |

Here's the problem. Given claim (iv), the gap in welfare for Amy (in red) provides a strong requiring reason to do (A) rather than (C). Any requiring reasons to do (C) rather than (A) are weaker, according to person-affecting views, so that there's an *overall* requiring reason to do (A) rather than (C). Similarly, the gap in welfare for Bobby (in blue) generates an overall requiring reason to do (B) rather than (A), and the gap in welfare for Carley (in green) generates an overall requiring reason to do (C) rather than (B). For each alternative, there's

²⁰ Also consider:

Same-Life Indifference: You have three alternatives. Create no one, create Amy with 60 years of happiness and no suffering, create Amy with 61 years of happiness and 1 year of suffering.

While the year of suffering generates a requiring reason against creating Amy with this suffering, the additional year of happiness is, according to claim (iv), a requiring reason to create her with 61 years of happiness *rather than* 60. If one year of happiness perfectly compensates for one year of suffering, then it is plausible that the requiring reason implied by (iv) balances the requiring reason not to create Amy with a year of suffering, so that there is no overall requiring reason to create Amy with one life rather than with the other. Moreover, there is a sufficiently strong permitting reason to create Amy with 61 years of happiness and 1 year of suffering rather than create no one. So, if it includes claim (iv), the Person-Affecting View has the plausible implication that in *Same-Life Indifference* it is permissible to create no one, and it is permissible to create Amy with either life.

²¹ Derek Parfit presented this sort of case to Larry Temkin in the 1980s. See Temkin (2012, 428-429).

an alternative that there is an overall requiring reason to do instead. This threatens to make each alternative wrong. But that would be implausible. Even if there are moral dilemmas, it seems implausible that this is one of them.

Here are three possible responses to this problem. First, we can deny that claim (iv) implies a requiring reason to do (A) over (C), a requiring reason to do (B) over (A), and a requiring reason to do (C) over (B). We could interpret (iv) to imply that, in such a three-alternative case, there is an equally strong requiring reason against each alternative. Second, we could allow that (iv) implies a cycle of “overall requiring reason to do _ rather than _” but maintain that the relevant requiring reasons are sufficiently balanced by permitting reasons, so that each alternative is permissible.²² Third, we can allow that there is a cycle of “overall requiring reason to do _ rather than _” but revise the Reasons-Wrongness Link, so that it’s no longer a tournament-style choice rule. Instead of saying that it’s permissible to choose an alternative only if it isn’t pairwise defeated, we can say that, when each alternative is pairwise defeated, it is permissible to choose one that is *defeated to the least degree*.²³

I won’t here take a stand on how the Person-Affecting View should be developed, to deal with the problem raised by *Amy, Bobby, and Carley*. Instead, my focus is the non-identity problem. I believe we can solve the non-identity problem somewhat independently of problems raised by cases like *Amy, Bobby, and Carley*.²⁴

5. Basic Non-Identity


So, let’s now return to the non-identity problem. In much of the literature on population ethics, the non-identity problem is discussed in terms of *overall welfare levels*, as follows.

²² This, roughly speaking, is the solution favored by Podgorski (2023) and Thomas (2023).

²³ For discussion of this kind of approach albeit in a different context, see Mann (2022).

²⁴ For related problems of “improvable life avoidance” see Ross (2015), Horton (2021a), Podgorski (2023), and Thornley (2023).

Overall Welfare Non-Identity: You have three alternatives. Create no one, create Amy at welfare 40, create Bobby at welfare 60. (Often this kind of case is presented using rectangles to represent overall welfare levels, as below.)




Many share the intuition that it is wrong to create Amy at 40, even if it is permissible to create no one at all. The problem is that the Person-Affecting View seems to be unable to capture this intuition, given that it says there's no requiring reason to create lives insofar as they contain happiness or positive welfare. Claim (iv) won't help here, since Amy and Bobby are different persons. (Similarly, the Person-Affecting View appears to be unable to capture the intuition that we should adopt a policy of reducing our carbon emissions in the *Carbon Policy* non-identity case I mentioned earlier.)

Most working in the area have sought to capture the “non-identity intuition” that it's wrong to create Amy in cases like *Overall Welfare Non-Identity*.²⁵ However, I believe that the common way of thinking about the non-identity problem, by focusing on overall welfare levels, overlooks important differences between the welfare components that make up these overall levels. I suggest that it matters whether Amy's overall welfare is lower than Bobby's because Amy's life contains more suffering than Bobby's, or merely because Amy's life contains less happiness than Bobby's.

²⁵ Exceptions include Roberts (2007), Boonin (2014), McDermott (2019), and Horton (2021a).

Let us first consider non-identity cases in which Amy's overall welfare is lower than Bobby's because her life contains more suffering. Let's start with this case:

Basic Non-Identity: You have three alternatives. Create no one, create Amy with 60 years of happiness and 6 months of suffering, create Bobby with 60 years of happiness and no suffering.



In this case, it seems wrong to create Amy, even though it is permissible to create no one. This intuitive claim is implied by claims (i)-(iii) together with the Reasons-Wrongness Link. First, note that (i) and (ii) imply that the only requiring reason present here is the requiring reason not to create Amy. There is therefore an overall requiring reason to create no one rather than create Amy, and an overall requiring reason to create Bobby rather than create Amy. According to the Reasons-Wrongness Link, then, it is permissible to create Amy *only if* there is a sufficiently strong permitting reason to create Amy rather than create no one *and* a sufficiently strong permitting reason to create Amy rather than Bobby.

While there is a sufficiently strong permitting reason to create Amy rather than create no one, there is *no* permitting reason to create Amy rather than Bobby. According to (iii), there is a permitting reason to create someone *rather than* create someone else or no one at all, insofar as the former alternative brings about *more* happiness than the latter alternative. The permitting reason here is contrastive. Since Amy's life would contain no more happiness

than Bobby's, (iii) would imply no permitting reason to create Amy rather than Bobby. If creating Bobby were not an alternative, it would be permissible to create Amy, since then there'd be a sufficiently strong permitting reason to create Amy rather than do the only alternative of creating no one. But, when creating Bobby is an alternative, there is no sufficiently strong permitting reason to create Amy rather than create Bobby.²⁶

This is how the Person-Affecting View solves the non-identity problem for cases in which Amy's overall welfare is lower than Bobby's because her life contains more suffering. Note that, while I claim that it is wrong to create Amy in *Basic Non-Identity*, I need not claim that doing so *wrongs* Amy. On one view, creating Amy does wrong her, since *her* pro tanto complaint is what grounds the fact that creating her is wrong: her pro tanto complaint generates a requiring reason not to create her and there is no sufficiently strong permitting reason to create her rather than Bobby. But, on another view, creating Amy does not wrong

²⁶ It may be important to observe here that there are crucial differences between *creating* new people and *saving* the lives of people who already exist. Consider the following case:

Saving Lives: You can at no cost to yourself or anyone else save the life of one of two people: Dylan or Ester. If you save Dylan, his future will contain 60 years of happiness and 6 months of suffering. If you save Ester, her future will contain 60 years of happiness and no suffering. You can save neither, save Dylan, or save Ester. You cannot save both.

There is a strong requiring reason to save each life. Since there is no permitting reason not to help, you are required to save someone. But who? Let us assume that the requiring reason to save Ester is stronger than the requiring reason to save Dylan, so that there is an overall requiring reason to save Ester rather than Dylan. The analogue of claim (iii) suggests that, since saving Dylan would bring about no more happiness than saving Ester, there is no permitting reason to save Dylan rather than Ester, so that you are required to save Ester. Many find this implication counterintuitive—indeed, Hanser (2009, 183) raises this kind of case as an objection to Harman's solution to the non-identity problem.


However, it seems to me that, since there is a strong requiring reason to save Dylan *in particular* and a strong requiring reason to save Ester *in particular*, there is a permitting reason to save *each* person in *Saving Lives*, reflecting the particularity of the requiring reasons in play. The permitting reason to save Dylan rather than Ester is strong enough to make it permissible to save Dylan, even if there's an overall requiring reason to save Ester rather than Dylan—in Pummer (2023, chapter 2), I defend the existence of such "individualist" permitting reasons. Alternatively, many argue that, in the face of a strong requiring reason to save Dylan in particular and a strong requiring reason to save Ester in particular, fairness requires you to give them each an equal chance of being saved, even if the requiring reason to save Ester is stronger than the requiring reason to save Dylan.

By contrast, in *Basic Non-Identity*, there is—according to my claim (ii) anyway—*no* requiring reason to create Amy in particular, and *no* requiring reason to create Bobby in particular, therefore there is neither a permitting reason to create each person in particular nor is there a requirement of fairness to give each an equal chance of being created. In short, we should not expect the particularity of happy people we can create to matter morally in the way the particularity of independently existing people matters morally.

her, given that she would have no overall complaint against being created. I do not need to take a stand either way. I wish only to note that, *if* the second of these views is correct, so that creating Amy does not wrong her, then we ought to reject person-affecting views according to which an act is wrong only if it *wrongs* someone in particular. On my Person-Affecting View, it is wrong to create Amy, even though she lacks an overall complaint against being created.

To further clarify how my version of the Person-Affecting View works, consider some variants of *Basic Non-Identity*.

Basic Non-Identity II: You have three alternatives. Create no one, create Amy with 60 years of happiness and 6 months of suffering, create Bobby with 60 years of happiness and 1 month of suffering.




Claims (i) and (ii) imply that there are only two requiring reasons in play: a requiring reason not to create Amy and a weaker requiring reason not to create Bobby. It is permissible to create no one, since this alternative is unopposed by requiring reasons.

What about creating Bobby? Since the requiring reason not to create Amy is significantly stronger than the requiring reason not to create Bobby, there is an overall requiring reason to create Bobby rather than create Amy (so, the alternative of creating Bobby isn't defeated by the alternative of creating Amy). There is an overall requiring reason to create no one rather than create Bobby, but there is a sufficiently strong permitting reason

to create Bobby rather than create no one (so, the alternative of creating Bobby isn't defeated by the alternative of creating no one). Therefore, by the Reasons-Wrongness Link, it is permissible to create Bobby.

What about creating Amy? There is an overall requiring reason to create no one rather than create Amy, but there is a sufficiently strong permitting reason to create Amy rather than create no one. However, there is an overall requiring reason to create Bobby rather than create Amy, and no permitting reason to create Amy rather than create Bobby. So, by the Reasons-Wrongness Link, it is wrong to create Amy. My Person-Affecting View treats *Basic Non-Identity II* much like *Basic Non-Identity*. Consider another variant.

Basic Non-Identity III: You have three alternatives. Create no one, create Amy with 60 years of happiness and 6 months of suffering, create Bobby with 40 years of happiness and 1 month of suffering.



This is just like *Basic Non-Identity II* with respect to requiring reasons. The key difference is that now there is a permitting reason to create Amy rather than create Bobby. Creating Amy brings about 5 additional months of suffering, but it also brings about 20 additional years of


happiness. I here leave it an open question whether this additional happiness provides a sufficiently strong permitting reason. If it does, then it is permissible to create Amy.²⁷

We have now seen how the Person-Affecting View solves the non-identity problem for cases like *Basic Non-Identity* in which Amy's overall welfare is lower than Bobby's because her life contains more suffering. The key to this solution is claim (i), which, recall, is supported by the thought that an individual's pro tanto complaints against being created can generate requiring reasons not to create them, even when they lack overall complaints against being created.

6. Non-Suffering Non-Identity

Now consider a case in which Amy's welfare is lower than Bobby's merely because her life contains less happiness than Bobby's.

Non-Suffering Non-Identity: You have three alternatives. Create no one, create Amy with 40 years of happiness and no suffering, create Bobby with 60 years of happiness and no suffering.



²⁷ We should be careful not to conclude that it's permissible to create Amy in this case on the basis of similar same-life additions: for example, the fact there's a sufficiently strong permitting reason to create Amy with 20 years of happiness and 5 months of suffering rather than create no one *does not clearly* imply that there's a sufficiently strong permitting reason to create Amy rather than Bobby in *Basic Non-Identity III*. The latter case involves additions of happiness and suffering *across different lives*: for example, the "additional 20 years of happiness" consists in the difference between Amy's 60 years and Bobby's 40 years. Having said this, we should allow that there is some difference in happiness between Amy and Bobby that would justify the difference in suffering between them so that it is permissible to create Amy.

On their own, claims (i)-(iv) imply that each alternative is permissible. This shouldn't come as much of a surprise. According to the Person-Affecting View, there is a requiring reason to do act *A* if and only if there would be something to be said against *not-A* from some particular individual's point of view. There is nothing to be said against any of these alternatives, from any particular individual's point of view. Not only does Amy lack any overall complaint against being created, she also lacks any pro tanto complaint against being created. Unlike in *Same-Life Improvement*, in *Non-Suffering Non-Identity*, there is nothing to be said against creating Amy with 40 years of happiness, from Amy's point of view: you would have made her as well off as you could have done.

Still, many report the intuition that it is wrong to create Amy in *Non-Suffering Non-Identity*. But I believe there is good reason to doubt the reliability of this intuition. I suspect many will have this intuition because they are antecedently committed to the "non-identity intuition" in cases like *Overall Welfare Non-Identity*. And it is unlikely that, when people imagine cases like *Overall Welfare Non-Identity*, they are imagining a case like *Non-Suffering Non-Identity*. Indeed, it is difficult to imagine *Non-Suffering Non-Identity* properly, so that there is absolutely no ill-being (whether hedonic or not) in Amy's life. It is very natural to imagine that Amy would, like nearly all of us, go through various trials and tribulations throughout her development as a person. But instead, we need to try and imagine something like this: if you push a button, Amy will pop into existence as a fully grown adult human being, on some distant uninhabited planet. Amy will take pleasure in food, natural beauty, and contemplating mathematics. She'll never encounter anything painful in her environment. Her psychology is unusual in that she'll never feel down, not even mildly bored. None of her projects or desires will be frustrated. Amy will be perfectly content with her relatively short lifespan. I suggest that, if we imagine *Non-Suffering Non-Identity*

properly so that it's sufficiently clear that Amy wouldn't have even a pro tanto complaint against being created, most of those independently attracted to the Procreation Asymmetry will lose the intuition that it is wrong to create her.²⁸

Now, some might suggest that Amy *would* have a pro tanto complaint, in virtue of not having *enough happiness*. Some might accept the following.

Shortfall View: Lives that contain under 50 years of happiness (positive welfare) fall short of a *standard of adequacy*. There is a requiring reason not to bring about a life that falls short in this way (because the person who would have such a life would have a pro tanto complaint against being created).²⁹

On this view, there is a requiring reason not to create Amy in *Non-Suffering Non-Identity*, just in virtue of the fact that her life falls short of the standard of adequate lifetime happiness (50 years of happiness). There is no such requiring reason not to create Bobby since his life exceeds the standard. Thus, if the Shortfall View of welfare is correct, then cases like *Non-Suffering Non-Identity* are relevantly like *Basic Non-Identity*, so that it is wrong to create the worse off person (Amy) in each case.

Having said all this, I am skeptical of the Shortfall View. The issue concerns the strength of the requiring reason not to create lives that fall below the standard of adequate lifetime happiness. Of course, the Shortfall View can say different things about *how strong* the requiring reason generated by falling short is. For example, it is compatible with the plausible claim that, in a choice between creating Amy with 40 years of happiness and no

²⁸ It may also be difficult to properly imagine that there is absolutely no suffering in *Same-Life Improvement*, but in that case, it seems clear that even if there were no suffering involved, Amy would have a pro tanto complaint against being created with less happiness *rather than* being created with more happiness.

²⁹ This is meant as an illustration of a broader family of views. Harman (2004) and Rivera-López (2009) set the standard in terms of normality, whereas Masny (2023) and Aaron (2024) set it in terms of potentiality.

suffering or creating no one at all, it is permissible to create Amy. But presumably the view must say that there is a *significant* requiring reason not to create a life that falls very far short of the standard (otherwise the view will not have significant implications). But then, as we consider possible lives with less and less happiness, at some point the Shortfall View will imply that the requiring reason generated just by falling short is stronger than the permitting reason to create the life, generated by the happiness that life would contain, so that it is wrong to create someone with this life, even if the only alternative is creating no one at all. For example, suppose you can either do nothing or open your fridge to get a slice of watermelon. Suppose that somehow opening the door would as a side effect press a button that would cause Amy to pop into existence in some remote uninhabited location as a fully grown adult. Amy would live for one moderately happy day. She would be perfectly content, and her life would contain nothing sad or boring. There would be no desire frustration. She would die peacefully watching the sunset. (If you like, you can imagine that Amy is an unusually short-lived goat.) The Shortfall View would presumably imply that it is wrong to create Amy. While it is difficult to properly imagine such cases, when I do my best to imagine this sort of case, it seems to me permissible to open my fridge. For these sorts of reasons, I am skeptical of the Shortfall View.³⁰

More generally, when I reflect on cases like *Non-Suffering Non-Identity*—which involve absolutely no ill-being whatsoever—along with the elegance and power of the Person-Affecting View as I’ve interpreted it, I find it plausible that you may create either person. At the very least, it is worth recognizing that, if you wish to capture the claim that it’s wrong to create Amy while retaining the Procreation Asymmetry, you will have to go beyond

³⁰ Similarly, I am skeptical of axiological views like Critical-Level Utilitarianism insofar as they imply that opening my fridge makes the world a worse place. For discussion of such views, see, for example, Broome (2004), Bykvist (2007), and Gustafsson (2020).

the simple Person-Affecting View: you'll have to posit some independent source of normativity. This will involve a loss of theoretical unity.

It has been suggested that we *can* square the claim that it is wrong to create Amy in *Non-Suffering Non-Identity* with the Procreation Asymmetry.³¹ While I cannot decisively show this here, this just strikes me as impossible. Attempts to do this end up blurring the distinction between non-identity cases like *Non-Suffering Non-Identity* and cases like *Same-Life Improvement*. In *Non-Suffering Non-Identity* we might be tempted to think there is something to be said against creating Amy, from the point of view of *the people you could create*—if you created Amy this would be worse for the collection of “Amy & Bobby” than if you created Bobby. But the *collection* of Amy & Bobby is not a particular individual, nor does it ground a particular individual's point of view.³² What makes it wrong to create Amy with 40 years of happiness in *Same-Life Improvement* is that there would be something to be said against this act from a particular individual's point of view. That simply fails to hold in *Non-Suffering Non-Identity*.³³


I have argued that, partly due to the difficulty of properly imagining the case, we can accept the implication that it is permissible to create Amy in *Non-Suffering Non-Identity*. But now consider:

³¹ For example, see Kumar (2003), Hare (2007), and Frick (2020).

³² For relevant discussion, see Boonin's (2014) (2019) and Harney's (2023) responses to views like Hare's (2007) and Kumar's (2003).

³³ For similar reasons, I am skeptical of Frick's (2020) attempt to square the Procreation Asymmetry with the claim that it is wrong to create Amy in cases like *Non-Suffering Non-Identity* (he actually focuses on a case like *Overall Welfare Non-Identity*, but what he says would carry over to the non-suffering variant). Frick appeals to a general “principle of standard selection” (79) which, applied to this case, says that since creating Amy satisfies the “Amy-standard” to a lesser degree than creating Bobby satisfies the “Bobby-standard,” and since these standards are of the same kind, there is a (requiring) reason to create Bobby *rather than* create Amy. First, it is not entirely clear to me how creating Amy satisfies the Amy-standard to a lesser degree than creating Bobby satisfies the Bobby-standard. But even assuming Frick is right about this, I find it implausible that these standards are of the same kind in the relevant sense that would support a requiring reason to create Bobby rather than Amy. At least, on a Person-Affecting View that is “narrow” with respect to wrong-makingness (like mine), there is a requiring reason to promote welfare if and only if there would be something to be said against failing to do so, from some particular individual's point of view. But there would be nothing to be said against creating Amy, from any particular individual's point of view. In claiming that the Amy-standard and the Bobby-standard are *of the same kind*, in the relevant sense that would support a requiring reason to create Bobby rather than Amy, Frick is going beyond this sort of narrow Person-Affecting View.

Equal Suffering Non-Identity: You have three alternatives. Create no one, create Amy with 40 years of happiness and 6 months of suffering, create Bobby with 60 years of happiness and 6 months of suffering.



To adequately grasp this case, we don't have to imagine that Amy's life contains absolutely no ill-being. Instead, we only have to imagine that the difference in welfare between Amy and Bobby is entirely a matter of a difference in happiness. It is plausible that's something we can imagine.

Yet, on the Person-Affecting View *as I have spelled it out thus far*, it's permissible to create Amy: there's no overall requiring reason to create Bobby rather than Amy, and while there's an overall requiring reason to create no one rather than create Amy, there's a sufficiently strong permitting reason to create Amy rather than create no one. The fact that Bobby's life would contain more happiness than Amy's does not provide a requiring reason to create Bobby rather than Amy.

Still, many share the intuition that it is wrong to create Amy in *Equal Suffering Non-Identity*. I believe this intuition can be vindicated by my Person-Affecting View when it is supplemented with a plausible general principle. I will explain how in the next section.

7. Silencing

In many of the cases we have looked at so far, the requiring reasons not to create happy people based on the suffering their lives would contain can weigh against countervailing requiring reasons. For example, in *Basic Non-Identity II*, the requiring reason not to create Amy weighs against the requiring reason not to create Bobby. In many other cases, however, it is plausible that such requiring reasons are (partly) *silenced*, so that they do not contribute (as much) to the overall balance of requiring reasons.³⁴ Silencing consists in the partial or complete reduction of a reason's strength; it is equivalent to what others have called *attenuation*.³⁵

To appreciate the importance of silencing, consider the following variant of *Happy Amy*:

Broken Toe: You can press a red button that will prevent an independently existing person, Fred, from breaking his toe, and as a side effect cause Amy to pop into existence. Her life would contain 60 years of happiness and 6 months of suffering.

I find it plausible that you are required to press the button. However, without *silencing*, my view would seem to imply that you are not required to press the button. According to the Reasons-Wrongness Link, it is wrong to do an act A if and only if there is an overall morally requiring reason to do an alternative B instead, and insufficient permitting strength supporting this act A over the alternative B. In the absence of any silencing phenomenon of the sort recognized by the Silencing Principle, permitting reasons merely contribute toward preventing the fact that there is an overall requiring reason to do B rather than A from making

³⁴ I am grateful to Tomi Francis for raising a question at the 2024 Parfit Lecture that helped me see this.

³⁵ See Dancy (2004) and Bader (2016).

A wrong. When the permitting reasons to do A are sufficiently strong, they successfully prevent the overall balance of requiring reasons from making A wrong. In the absence of any silencing, permitting reasons do not affect the strength of requiring reasons and thus do not affect how requiring reasons contribute to the overall balance of requiring reasons. By contrast, the sort of silencing justification recognized by the Silencing Principle would instead partly or completely reduce the strength of requiring reasons, affecting how they contribute to the overall balance of requiring reasons. Silencing in this sense vindicates the intuition that you are *required* to press the button in *Broken Toe*. If Amy's happiness did *not* affect the strength of the requiring reason not to create Amy, based on her suffering, then it would presumably remain the case that there is an overall requiring reason not to press the button rather than press the button (as presumably the requiring reason not to cause Amy 6 months of suffering is stronger than the requiring reason to prevent Fred from breaking his toe). So, according to the Reasons-Wrongness Link, it would remain permissible not to press the button. To vindicate the intuition that you are required to press the button—that it is wrong not to press the button—we need it to be the case that there is an overall requiring reason to press the button rather than not. This is what the Silencing Principle delivers: with the requiring reason not to create Amy silenced, it doesn't contribute to the overall balance of requiring reasons, so that the requiring reason to prevent Fred from breaking his toe makes it the case that there's an overall requiring reason to press the button rather than not. From the Reasons-Wrongness Link together with fact that there's no permitting reason not to press the button, it follows that you are required to press the button.

In sum, I suggest that what explains the intuition that you are required to press the button in *Broken Toe* is the idea that, if Amy's suffering is *justified in the right way*, it has been “dealt with” and therefore fails to exert (significant) moral pressure that could clash

with other moral pressures, such as that to prevent Fred from breaking his toe. Here is a principle that captures the claim that you are required to press the button in *Broken Toe*.

Silencing Principle: To the extent that *A*'s happiness provides permitting strength supporting creating *A* rather than doing an alternative, the requiring reason to do this alternative rather than create *A* (based on *A*'s suffering) is silenced.³⁶

While this principle appeals to permitting reasons, the sort of justificatory role it recognizes is distinct from that of permitting reasons. Again, permitting reasons on their own contribute toward making an act permissible *in the face of* countervailing requiring reasons, without affecting the strength of these requiring reasons. The sort of silencing justification recognized by the Silencing Principle would instead partly or completely reduce the strength of requiring reasons. It is plausible that such silencing happens in *Broken Toe*.

Now consider the following case.

Broken Toe Non-Identity: You have three alternatives. Create no one, press a red button that prevents Fred from breaking his toe and as a side effect creates Amy with 60 years of happiness and 6 months of suffering, or press a green button that creates Bobby with 60 years of happiness and no suffering.

In this case, it seems to me wrong to press the red button. In this case, Amy's suffering is not "justified in the right way." While Amy's suffering is *intrapersonally compensated* by her own happiness, such compensation is not sufficient to silence her suffering (although such


³⁶ It is plausible that silencing is mutual: to the extent that the permitting strength supporting creating *A* rather than doing an alternative silences the requiring reason to do this alternative rather than create *A*, this permitting strength is also silenced; these countervailing permitting and requiring forces in this way neutralize each other.

intrapersonal compensation may be necessary for silencing). As implied by claim (iii), the happiness in Amy's life provides *no* permitting reason to press the red button rather than the green button, given that the latter would create Bobby with *no less* happiness. The Silencing Principle leaves unsilenced the requiring reason to press the green button rather than the red button, based on the 6 months of suffering Amy's life would contain. This unsilenced requiring reason is intuitively stronger than the requiring reason to prevent Fred from breaking his toe, so that there is an overall requiring reason to press the green button rather than the red button. Given this, and the fact that there's no sufficiently strong permitting reason to press the red button rather than the green button, it is wrong to press the red button.³⁷

The Silencing Principle has the intuitively plausible implications that you are required to press the red button in *Broken Toe*, and that it is wrong to press the red button in *Broken Toe Non-Identity*. The Silencing Principle has plausible implications in other cases, too. Recall:

³⁷ Bonus question: is it also wrong to create no one? The Silencing Principle implies that, since Amy's happiness *does* provide a permitting reason to create Amy rather than create no one, the requiring reason to create no one rather than create Amy is silenced. Given the requiring reason to prevent Fred from breaking a toe, there is an overall requiring reason to press the red button rather than create no one. Since there is no permitting reason to create no one rather than press the red button, it is wrong to create no one. Thus, you are required to press the green button, creating Bobby. This strikes me as a plausible result—you are required to press the red button *rather than* create no one and you are required to press the green button *rather than* the red button, leaving you with just the one undefeated alternative of pressing the green button. Having said this, note that it is possible to revise the Silencing Principle so that, for the requiring reason to create no one rather than Amy to be silenced by Amy's happiness, her happiness would have to generate a sufficiently strong permitting reason to create her rather than do any other alternative, including creating Bobby. This revised principle could then imply that the requiring reason to create no one rather than press the red button is unsilenced, so that there is an overall requiring reason to create no one rather than press the red button. This could in turn get the implication that it is permissible to create no one, wrong to press the red button, and permissible to press the green button. The choice between this revised Silencing Principle and the Silencing Principle as it appears in the main text above does not make a relevant difference to the remaining issues to be discussed here, so for the sake of simplicity I will continue to work with the Silencing Principle as it appears in the main text.

Equal Suffering Non-Identity: You have three alternatives. Create no one, create Amy with 40 years of happiness and 6 months of suffering, create Bobby with 60 years of happiness and 6 months of suffering.



The Silencing Principle can explain why it is wrong to create Amy in this case. First, recall that, according to claim (i), there is a requiring reason not to create Amy and there is a requiring reason not to create Bobby. Since neither reason is stronger than the other, there is no *overall* requiring reason to create one person rather than the other. Nonetheless, there is *a* requiring reason to create Amy rather than Bobby (and *a* requiring reason to create Bobby rather than Amy). Next, notice that according to claim (iii), given that Bobby would have 20 more years of happiness than Amy, there is a correspondingly strong permitting reason to create Bobby rather than Amy. Since this permitting reason is based on (these 20 years of) *Bobby's happiness*, the Silencing Principle implies that the requiring reason to create Amy rather than create Bobby (based on Bobby's suffering) is *to some considerable extent* silenced. Since the remaining unsilenced requiring reason to create Amy rather than Bobby (based on Bobby's suffering) is therefore to a considerable extent weaker than the requiring reason to create Bobby rather than Amy (based on Amy's suffering), there is an *overall* requiring reason to create Bobby rather than Amy. Since there is no permitting reason to create Amy rather than Bobby, it is wrong to create Amy.

For another success story for the Silencing Principle, consider:

Miserable Non-Identity: You have two alternatives. Create Amy with no happiness and 60 years of suffering or create Bobby with 50 years of happiness and 60 years of suffering.³⁸

It seems clear you are required to create Bobby. The Silencing Principle can explain this. According to claim (iii), given that Bobby would have 50 more years of happiness than Amy, there is a correspondingly strong permitting reason to create Bobby rather than Amy. Since this permitting reason is based on Bobby's happiness, the Silencing Principle implies that the requiring reason to create Amy rather than create Bobby (based on Bobby's suffering) is to some considerable extent silenced. As before, we get the result that it is wrong to create Amy.

Finally, consider the following case.

Miserable Non-Identity II: You have two alternatives. Press a red button that creates Amy with no happiness and 6 months of suffering *or* press a green button that creates 100 happy people each with 100 years of happiness and no suffering *and* as a side effect creates Bobby with no happiness and 3 years of suffering.

Let us assume, as seems plausible enough, that the happiness of the 100 happy people makes it permissible to press the green button, even though this involves creating Bobby as a side effect, bringing about more suffering than creating Amy does. Crucially, it is not *Bobby's* happiness that provides a permitting reason to press the green button, but the happiness of the

³⁸ If you like, you can imagine that you also have the alternative of creating no one. But we can suppose this third alternative would be catastrophic for very many independently existing people, so that you are required to choose of the two "creation" alternatives instead.

100 happy people. There is no intrapersonal compensation of Bobby's suffering. In this case, the Silencing Principle does *not* imply that the requiring reason to press the red button rather than the green button (based on Bobby's suffering) is silenced. And that seems to me the correct result. The unsilenced requiring reason to press the red button rather than the green button (based on Bobby's suffering) is stronger than the requiring reason to press the green button rather than the red button (based on Amy's suffering). Thus, there is an overall requiring reason to press the red button rather than the green button. There is also a sufficiently strong permitting reason to press the green button rather than the red button. This gives what seems to me the intuitively correct result that, while you are *permitted* to press the green button (creating the 100 happy people along with Bobby), you are not *required* to do so. Even if you do not agree that *Miserable Non-Identity II* is itself such a case, it does seem quite plausible that there will be many cases in which suffering is justified in a way that *makes it permissible* to create lives that contain this suffering, but is *not* justified in a way that "deals with" the suffering so that the requiring reasons not to create these lives are *silenced*.

Cases like *Broken Toe* and others discussed in this section suggest that requiring reasons not to create people based on the suffering their lives would contain can be (at least to some considerable extent) silenced. The Person-Affecting View should include *some* such silencing role, but I am not wedded to the Silencing Principle in particular. I think it is largely an open question for further investigation, exactly how silencing should operate. Indeed, I also leave it largely open how to *weigh* requiring reasons to benefit existing people against the requiring reasons and permitting reasons generated by the Person-Affecting View, as specified by claims (i)-(iv).³⁹

³⁹ For relevant discussion of this kind of weighing issue, see Woollard (2012).

8. Practical Implications

Many of the non-identity cases I have focused on so far have the following sort of form: create no one, create A, or create B. In more realistic non-identity cases, our alternatives will involve creating multiple different people. Recall *Carbon Policy* (the more “real-life” non-identity case from the beginning of the paper). In *Carbon Policy*, there are three relevant groups of people: the 8 billion of us, the trillions of less well-off people who would exist in two centuries if we carry on with business as usual, and the trillions of better-off people who would exist in two centuries if we adopt the carbon policy. How my Person-Affecting View applies to this sort of case depends on the details. For example, it seems likely that in real-life versions of *Carbon Policy*, the less well-off people would have lower welfare at least partly due to having more suffering in their lives. If *Carbon Policy* is then relevantly like a scaled-up version of *Basic Non-Identity* (or *Basic Non-Identity II*), my view straightforwardly implies that it is wrong to carry on with business as usual. In further versions of *Carbon Policy*, the less well-off future people and better-off future people will make up groups of very *different sizes*, with very different welfare profiles. As my brief remarks about aggregation suggest, it may not be obvious how best to extend the Person-Affecting View to some of these cases. But it is plausible that it implies there is moral pressure to adopt the carbon policy in a wide range of these cases.


Choices between large-scale policies in global health affect the quality of life of future people. Our choice of policy will also almost always be identity-affecting, in the way it is in *Carbon Policy*. Consider the following non-identity case.

Life-Saving Vaccines: We can invest in vaccines that will prevent future deaths. If we don’t invest, many future people will die at the age of 40. If we do invest, *these*

particular future people will never exist, and the same number of different future people will die at the age of 60.

My view implies there are suffering-based requiring reasons to invest in these vaccines. In real-world contexts, disease-induced deaths cause many forms of suffering to those who die as well as to those around them. But even in many cases in which these deaths do not cause any suffering, my view still has the implication that there is an overall requiring reason to invest in the vaccines rather than not. This will be so if *Life-Saving Vaccines* is like a scaled-up version of *Equal Suffering Non-Identity*, so that there is equal suffering in the lives of each group of future people and more happiness in the lives of those who live longer. Of course, in realistic cases, longer lives don't just contain more happiness, they also contain more suffering. It is more plausible to regard *Life-Saving Vaccines* as a scaled-up version of this sort of case:

Proportionate Suffering Non-Identity: You have three alternatives. Create no one, create Amy with 40 years of happiness and 6 months of suffering, create Bobby with 60 years of happiness and 9 months of suffering. (An additional 20 years of happiness comes with 3 more months of suffering.)



Since the permitting reason to create Bobby rather than Amy is based on (these 20 years of) Bobby's happiness, the Silencing Principle implies that the requiring reason to create Amy rather than create Bobby (based on Bobby's suffering) is to some considerable extent silenced. If the "rate of silencing" is such that Bobby's 20 years of happiness silences *more* than 3 months of suffering, then there is an overall requiring reason to create Bobby rather than Amy. Since there is no permitting reason to create Amy rather than Bobby, it is wrong to create Amy. Therefore, to the extent that *Life-Saving Vaccines* is relevantly like a scaled-up version of *Proportionate Suffering Non-Identity*, my view, supplemented with something like the Silencing Principle, will imply that the suffering of the people living to 60 is silenced to a considerable extent, so that there is, overall, significant moral pressure to invest in the life-saving vaccines, bringing about the future people living to 60 rather than the different future people living to 40. And it is plausible that realistic versions of *Life-Saving Vaccines* will be relevantly like scaled-up versions of *Proportionate Suffering Non-Identity*. In gloomier versions of *Life-Saving Vaccines* in which additional years bring considerable additional suffering and not enough additional happiness, my view has the plausible implication that it is permissible (or even required) not to invest in the vaccines.

I have just argued that my view has plausible practical implications in more realistic cases like *Carbon Policy* and *Life-Saving Vaccines*. Consider another case:

Future People Lottery: We can use a fixed pot of resources to save the lives of 100 independently existing people for certain, or instead bring about a one-in-a-billion chance of creating 10^{20} happy people in the far future (or prevent an event that would prevent their existence).

One might worry that my view would have the counterintuitive implication that it is *permissible* to abandon the 100 independently existing people for the sake of playing this future people lottery. But whether my view has this implication depends on what theory of aggregation and risk it adopts. In discussing *Expectably Happy Amy*, I said that the chance of creating happiness provides a permitting reason to create Amy. Presumably the same would hold in *Future People Lottery*: the one-in-a-billion chance of creating the 10^{20} happy people provides a permitting reason to run this lottery. However, I need not commit to the view that the *strength* of this permitting reason is given by anything like multiplying this chance by the total amount of happiness that would be produced if the lottery is won. As the literature on partial aggregation and risk suggests, it is difficult to know what to say about these matters when dealing with requirements to save the lives of independently existing people (must we save 100 lives for sure rather than bring about a one-in-a-million chance of saving 200 million lives?).⁴⁰ Similar difficulties, and perhaps some additional ones, will appear when dealing with aggregation, risk, and permitting reasons to create new people.

As I suggested above in the context of Different-Lives Antinatalism, I do find it plausible that the numbers matter: there is a stronger permitting reason to create *more* happy people (as there would be more to be said in favor of this act, considering the points of view of those who would get to exist). If it were *very likely* that performing some act would result in the creation of 10^{20} happy people in the future, I find it plausible that it would be permissible to perform this act even if doing so meant failing to save the lives of 100 independently existing people. But such claims are not essential to the Person-Affecting View I wish to put forward here. I leave it an open question what theory of aggregation for permitting reasons we should combine with this view.

⁴⁰ See, for example: Frick (2015), Horton (2020), Walen (2020). On aggregation, risk, and future people, see Curran (2022) and Heikkinen (2022).

9. Conclusion

I've presented a new version of the Person-Affecting View, according to which reasons to do an act depend wholly on what would be said for or against this act from the points of view of particular individuals. According to my view, (i) there is a morally *requiring* reason not to bring about lives insofar as they contain suffering (negative welfare), (ii) there is *no* morally requiring reason to bring about lives insofar as they contain happiness (positive welfare), but (iii) there is a *permitting* reason to bring about lives insofar as they contain happiness (and no permitting reason to create one rather than the other, if the first contains no more happiness than the second). In a wide range of cases, my view can capture the "non-identity intuition" that it is wrong to create less well-off future people when we can instead create different better-off future people. It can do so while retaining the Procreation Asymmetry and avoiding implausible forms of antinatalism.

There is powerful moral pressure to improve the future. For example, my view implies there is such pressure to tackle climate change and avoid various forms of civilizational collapse. However, my view recognizes no moral pressure to colonize space if the only upshot of doing so is increasing total happiness. Many believe that the vast potential happiness of future people itself provides overwhelming moral pressure to avoid extinction. On my view, this happiness provides no such pressure. At the same time, on my view, future happiness is a source of permitting reasons. There'd be something to be said in favor of bringing about happy futures, from the points of view of the particular individuals who'd inhabit them.

References

- Aaron, Jonas H. 2024. "A Less Bad Theory of the Procreation Asymmetry and the Non-Identity Problem." *Utilitas* 36: 35–49.
- Algander, Per. 2012. "A Defence of the Asymmetry in Population Ethics." *Res Publica* 18:145–157.
- Bader, Ralf. 2016. "Conditions, Modifiers, and Holism." In *Weighing Reasons*, eds. Errol Lord and Barry Maguire, Oxford: Oxford University Press, pp. 27-55.
- Beckstead, Nick. 2019. "A Brief Argument for the Overwhelming Importance of Shaping the Far Future." In *Effective Altruism: Philosophical Issues*, eds. Hilary Greaves and Theron Pummer, Oxford: Oxford University Press, pp. 80–98.
- Benatar, David. 2006. *Better Never to Have Been: The Harm of Coming into Existence*. Oxford: Oxford University Press.
- Boonin, David. 2014. *The Non-Identity Problem and the Ethics of Future People*. Oxford: Oxford University Press.
- Boonin, David. 2019. "Solving the Non-Identity Problem: A Reply to Gardner, Kumar, Malek, Mulgan, Roberts and Wasserman." *Law, Ethics and Philosophy* 7: 127-156.
- Bostrom, Nick. 2013. "Existential Risk Prevention as Global Priority." *Global Policy* 4: 15–31.
- Broome, John. 2004. *Weighing Lives*. Oxford: Oxford University Press.
- Bykvist, Krister. 2007. "The good, the bad, and the ethically neutral." *Economics and Philosophy* 23: 97–105.
- Chappell, Richard Yetter. 2017. "Rethinking the Asymmetry." *Canadian Journal of Philosophy* 47: 167-177.
- Curran, Emma. 2022. "Longtermism, Aggregation, and Catastrophic Risk." GPI Working Paper No. 18-2022.
- Dancy, Jonathan. 2004. *Ethics without Principles*. Oxford: Oxford University Press.
- Dancy, Jonathan. 2021. *Practical Thought: Essays on Reason, Intuition, and Action*. Oxford: Oxford University Press.
- Frick, Johann. 2015. "Contractualism and Social Risk." *Philosophy and Public Affairs* 43: 175–223.
- Frick, Johann. 2020. "Conditional reasons and the procreation asymmetry." *Philosophical Perspectives* 34: 53–87.

- Gardner, Molly. 2015. "A harm-based solution to the non-identity problem." *Ergo* 2: 427–444.
- Gert, Joshua. 2003. "Requiring and justifying: two dimensions of normative strength." *Erkenntnis* 59: 5–36.
- Glover, Jonathan. 1977. *Causing Death and Saving Lives*. Harmondsworth: Penguin Books.
- Gustafsson, Johan. 2020. "Population axiology and the possibility of a fourth category of absolute value." *Economics and Philosophy* 36: 81–110.
- Hanser, Matthew. 2009. "Harming and procreating." In *Harming Future Persons: Ethics, Genetics and the Nonidentity Problem*, eds. M. Roberts and D. Wasserman, 179–200. Dordrecht: Springer.
- Hare, Caspar. 2007. "Voices from Another World: Must We Respect the Interests of People Who Do Not, and Will Never, Exist?" *Ethics* 117: 498–523.
- Harman, Elizabeth. 2004. "Can we harm and benefit in creating?" *Philosophical Perspectives* 18: 89–113.
- Harman, Elizabeth. 2009. "Harming as causing harm." In *Harming Future Persons: Ethics, Genetics and the Nonidentity Problem*, eds. M. Roberts and D. Wasserman, 137–154. Dordrecht: Springer.
- Harney, Jonas. 2023. "The Interpersonal Comparative View of Welfare: Its Merits and Flaws." *The Journal of Ethics* 27: 369–91.
- Heikkinen, Karri. 2022. "Strong Longtermism and the Challenge from Anti-Aggregative Moral Views." GPI Working Paper No. 5-2022.
- Horgan, Terry, and Mark Timmons. 2010. "Untying a Knot from the Inside Out: Reflections on the 'Paradox' of Supererogation." *Social Philosophy and Policy* 27: 29–63.
- Horton, Joe. 2017. "The All or Nothing Problem." *Journal of Philosophy* 114: 94–104.
- Horton, Joe. 2020. "Aggregation, Risk, and Reductio." *Ethics* 130: 514–529.
- Horton, Joe. 2021a. "New and Improvable Lives." *Journal of Philosophy* 118: 486–503.
- Horton, Joe. 2021b. "Partial Aggregation in Ethics." *Philosophy Compass* 16: 1–12.
- Kamm, Frances. 2007. *Intricate Ethics: Rights, Responsibilities, and Permissible Harm*. New York: Oxford University Press.
- Kumar, Rahul. 2003. "Who Can Be Wronged?" *Philosophy and Public Affairs* 31: 99–118.
- Little, Margaret and Coleen Macnamara. 2017 "For Better or Worse: Commendatory Reasons and Latitude." *Oxford Studies in Normative Ethics* Vol. 7, ed. Mark Timmons, 138–160. Oxford: Oxford University Press.

- Mann, Kirsten. 2022. "Relevance and Nonbinary Choices." *Ethics* 132: 382–413.
- Masny, Michal. 2023. "Wasted Potential: The Value of a Life and the Significance of What Could Have Been." *Philosophy and Public Affairs* 51: 6–32.
- McDermott, Michael. 2019. "Harms and Objections." *Analysis* 79: 436–48.
- McMahan, Jeff. 1981. "Problems of population theory." *Ethics* 92: 96–127.
- Mogensen, Andreas. 2019. "Staking Our Future: Deontic Longtermism and the Non-Identity Problem." GPI Working Paper - No. 9-2019.
- Muñoz, Daniel. 2021. "Three Paradoxes of Supererogation." *Nous* 55: 699–716.
- Narveson, Jan. 1973. "Moral Problems of Population." *The Monist* 57: 62–86.
- Parfit, Derek. 1984. *Reasons and Persons*. Oxford: Oxford University Press.
- Parfit, Derek. 2011. *On What Matters, Volume Two*. Oxford: Oxford University Press.
- Podgorski, Abelard. 2023. "Complaints and Tournament Population Ethics." *Philosophy and Phenomenological Research* 106: 344–67.
- Portmore, Douglas. 2011. *Commonsense Consequentialism*. New York: Oxford University Press.
- Pummer, Theron. 2023. *The Rules of Rescue*. New York: Oxford University Press.
- Rivera-López, Eduardo. 2009. "Individual Procreative Responsibility and the Non-Identity Problem." *Pacific Philosophical Quarterly* 90: 99–118.
- Roberts, Melinda. 2007. "The non-identity fallacy: Harm, probability and another look at Parfit's depletion example." *Utilitas* 19: 267–311.
- Roberts, Melinda. 2011. "An Asymmetry in the Ethics of Procreation." *Philosophy Compass* 6: 765–776.
- Ross, Jacob. (2015). "Rethinking the person-affecting principle." *Journal of Moral Philosophy* 12: 428–61.
- Rulli, Tina. 2020. "Conditional Obligations." *Social Theory and Practice* 46: 365–390.
- Scanlon, Thomas M. 1998. *What We Owe to Each Other*. Cambridge: Harvard University Press.
- Shiffrin, Seana. 1999. "Wrongful life, procreative responsibility, and the significance of harm." *Legal Theory* 5: 117–148.

- Sinclair, Thomas. 2018. "Are We Conditionally Obligated to Be Effective Altruists?" *Philosophy and Public Affairs* 46: 36–59.
- Snedegar, Justin. 2017. *Contrastive Reasons*. Oxford: Oxford University Press.
- Spencer, Jack. 2021. "The procreative asymmetry and the impossibility of elusive permission." *Philosophical Studies* 178: 3819–3842.
- Temkin, Larry. 2012. *Rethinking the Good*. Oxford: Oxford University Press.
- Thomas, Teruji. 2023. "The asymmetry, uncertainty, and the long term." *Philosophy and Phenomenological Research* 107: 470–500.
- Thornley, Elliott. 2023. "The Procreation Asymmetry, Improvable-Life Avoidance, and Impairable-Life Acceptance." *Analysis* 83: 517–26.
- Tucker, Chris. Forthcoming. "The All or Nothing Ranking Reversal and the Unity of Morality." *Oxford Studies in Normative Ethics*.
- Voorhoeve, Alex. 2014. "How Should We Aggregate Competing Claims?" *Ethics* 125: 64–87.
- Walen, Alec. 2020. "Risks and Weak Aggregation: Why Different Models of Risk Suit Different Types of Cases." *Ethics* 131: 62–86.
- Woollard, Fiona. 2012. "Have We Solved the Non-Identity Problem?" *Ethical Theory and Moral Practice* 15: 677–690.
- Woollard, Fiona. 2015. *Doing and Allowing Harm*. Oxford: Oxford University Press.