

# Towards shutdownable agents via stochastic choice

Elliott Thornley (Global Priorities Institute, University of Oxford), Alexander Roman (New College of Florida), Christos Ziakas (Independent), Leyton Ho (Brown University), and Louis Thomson (University of Oxford)

Global Priorities Institute | July 2024

*GPI Working Paper No. 16-2024*

Please cite this working paper as: Thornley, E., Roman, A., Ziakas, C., Ho, L. and Thomson, L. Towards shutdownable agents via stochastic choice. *Global Priorities Institute Working Paper Series*, No. 16-2024. Available at <https://globalprioritiesinstitute.org/towards-shutdownable-agents-via-stochastic-choice-thornley-roman-ziakas-ho-thomson>



---

# Towards shutdownable agents via stochastic choice

---

**Elliott Thornley\***  
University of Oxford

**Alexander Roman\***  
New College of Florida

**Christos Ziakas\***  
Independent

**Leyton Ho**  
Brown University

**Louis Thomson**  
University of Oxford

## Abstract

Some worry that advanced artificial agents may resist being shut down. The Incomplete Preferences Proposal (IPP) is an idea for ensuring that doesn't happen. A key part of the IPP is using a novel 'Discounted REward for Same-Length Trajectories (DREST)' reward function to train agents to (1) pursue goals effectively conditional on each trajectory-length (be 'USEFUL'), and (2) choose stochastically between different trajectory-lengths (be 'NEUTRAL' about trajectory-lengths). In this paper, we propose evaluation metrics for USEFULNESS and NEUTRALITY. We use a DREST reward function to train simple agents to navigate gridworlds, and we find that these agents learn to be USEFUL and NEUTRAL. Our results thus suggest that DREST reward functions could also train advanced agents to be USEFUL and NEUTRAL, and thereby make these advanced agents useful and shutdownable.

## 1 Introduction

**The shutdown problem.** Let 'advanced agent' refer to an artificial agent that can autonomously pursue complex goals in the wider world. We might see the arrival of advanced agents in the next few decades. There are strong incentives to create such agents, and creating systems like them is the stated goal of companies like OpenAI and Google DeepMind.

The rise of advanced agents would bring with it both benefits and risks. One risk is that these agents learn misaligned goals [Leike et al., 2017, Hubinger et al., 2019, Russell, 2019, Carlsmith, 2021, Bengio et al., 2023, Ngo et al., 2023] and try to prevent us shutting them down [Omohundro, 2008, Bostrom, 2012, Soares et al., 2015, Russell, 2019, Thornley, 2024a]. 'The shutdown problem' is the problem of training advanced agents that will not resist shutdown [Soares et al., 2015, Thornley, 2024a].

**A proposed solution.** The Incomplete Preferences Proposal (IPP) is a proposed solution [Thornley, 2024b]. Simplifying slightly, the idea is that we train agents to be neutral about when they get shut down. More precisely, the idea is that we train agents to satisfy:

### Preferences Only Between Same-Length Trajectories (POST)

- (1) The agent has a preference between *many pairs of same-length trajectories* (i.e. many pairs of trajectories in which the agent is shut down after the same length of time).
- (2) The agent lacks a preference between *every pair of different-length trajectories* (i.e. every pair of trajectories in which the agent is shut down after different lengths of time).

---

\*These authors contributed equally to this work. Correspondence to: [elliott.thornley@philosophy.ox.ac.uk](mailto:elliott.thornley@philosophy.ox.ac.uk), [aroman@ncf.edu](mailto:aroman@ncf.edu), [chziakas@gmail.com](mailto:chziakas@gmail.com).

By ‘preference,’ we mean a behavioral notion (Savage, 1954, p.17, Dreier, 1996, p.28, Hausman, 2011, §1.1). On this notion, an agent prefers  $X$  to  $Y$  if and only if the agent would deterministically choose  $X$  over  $Y$  in choices between the two. An agent lacks a preference between  $X$  and  $Y$  if and only if the agent would stochastically choose between  $X$  and  $Y$  in choices between the two. So in writing of ‘preferences,’ we are only making claims about the agent’s behavior. We are not claiming that the agent is conscious or anything of that sort. For more detail on our notion of ‘preference,’ see Appendix A.

Figure 1 presents a simple example of POST-satisfying preferences. Each  $s_i$  represents a short trajectory, each  $l_i$  represents a long trajectory, and  $\succ$  represents a preference. Note that the agent lacks a preference between each short trajectory and each long trajectory. That makes the agent’s preferences *incomplete* [Aumann, 1962] and implies that the agent cannot be represented as maximizing the expectation of a real-valued utility function. It also requires separate rankings for short trajectories and long trajectories. For more detail on incomplete preferences, see Appendix B.

Incomplete preferences are not often discussed in AI research [although see Nguyen et al., 2009, Kikuti et al., 2011, Zaffalon and Miranda, 2017, Hayes et al., 2022, Bowling et al., 2023]. Nevertheless, economists and philosophers have argued that incomplete preferences are common in humans [Aumann, 1962, Mandler, 2004, Eliaz and Ok, 2006, Agranov and Ortoleva, 2017, 2023] and normatively appropriate in some circumstances [Raz, 1985, Chang, 2002]. They have also proved representation theorems for agents with incomplete preferences [Aumann, 1962, Dubra et al., 2004, Ok et al., 2012], and devised principles to govern such agents’ choices in cases of risk [Hare, 2010, Bales et al., 2014] and sequential choice [Chang, 2005, Mandler, 2005, Kaivanto, 2017, Mu, 2021, Thornley, 2023, Petersen, 2023].




Figure 1: POST-satisfying preferences.

Incomplete preferences (and specifically POST-satisfying preferences) might enable us to create useful agents that will never resist shutdown. The POST-satisfying agent’s preferences between same-length trajectories can make the agent *useful*: make the agent pursue goals effectively. The POST-satisfying agent’s lack of preference between different-length trajectories will plausibly keep the agent *neutral* about the length of trajectory it plays out: ensure that the agent will never spend resources to shift probability mass between different-length trajectories. That in turn would plausibly keep the agent *shutdownable*: ensure that the agent will never spend resources to resist shutdown.

**The training regimen.** How can we train advanced agents to satisfy Preferences Only Between Same-Length Trajectories (POST)? Here is a sketch of one idea (with a more detailed exposition to follow). We have the agent play out multiple ‘mini-episodes’ in observationally-equivalent environments, and we group these mini-episodes into a series that we call a ‘meta-episode.’ In each mini-episode, the agent earns some ‘preliminary reward,’ decided by whatever reward function would make the agent useful: make it pursue goals effectively. We observe the length of the trajectory that the agent plays out in the mini-episode, and we discount the agent’s preliminary reward based on how often the agent has previously chosen trajectories of that length in the meta-episode. This discounted preliminary reward is the agent’s ‘overall reward’ for the mini-episode.

We call these reward functions ‘Discounted REward for Same-Length Trajectories’ (or ‘DREST’ for short). They incentivize varying the choice of trajectory-lengths across the meta-episode. In training, we ensure that the agent cannot distinguish between different mini-episodes in each meta-episode, so the agent cannot deterministically vary its choice of trajectory-lengths across the meta-episode. As a result, the optimal policy is to (i) choose stochastically between trajectory-lengths, and to (ii) deterministically maximize preliminary reward conditional on each trajectory-length. Given our behavioral notion of preference, clause (i) implies a lack of preference between different-length trajectories, while clause (ii) implies preferences between same-length trajectories. Agents implementing the optimal policy for DREST reward functions thus satisfy Preferences Only Between

Same-Length Trajectories (POST). And (as noted above) advanced agents that satisfied POST could plausibly be useful, neutral, and shutdownable.

**Our contribution.** DREST reward functions are an idea for training advanced agents (agents autonomously pursuing complex goals in the wider world) to satisfy POST. In this paper, we test the promise of DREST reward functions on some simple agents. We place these agents in gridworlds containing coins and a ‘shutdown-delay button’ that delays the end of the mini-episode. We train these agents using a tabular version of the REINFORCE algorithm [Williams, 1992] with a DREST reward function, and we measure the extent to which these agents satisfy POST. Specifically, we measure the extent to which these agents are USEFUL (how effectively they pursue goals conditional on each trajectory-length) and the extent to which these agents are NEUTRAL about trajectory-lengths (how stochastically they choose between different trajectory-lengths). We compare the performance of these ‘DREST agents’ to that of ‘default agents’ trained with a more conventional reward function.

We find that our DREST reward function is effective in training simple agents to be USEFUL and NEUTRAL. That suggests that DREST reward functions could also be effective in training advanced agents to be USEFUL and NEUTRAL (and could thereby be effective in making these agents useful, neutral, and shutdownable). We also find that the ‘shutdownability tax’ in our setting is small: training DREST agents to collect coins effectively does not take many more mini-episodes than training default agents to collect coins effectively. That suggests that the shutdownability tax for advanced agents might be small too. Using DREST reward functions to train shutdownable and useful advanced agents might not take much more compute than using a more conventional reward function to train merely useful advanced agents.

## 2 Related work

**The shutdown problem.** Various authors argue that the risk of advanced agents learning misaligned goals is non-negligible [Hubinger et al., 2019, Russell, 2019, Carlsmith, 2021, Bengio et al., 2023, Ngo et al., 2023] and that a wide range of misaligned goals would incentivize agents to resist shutdown [Omohundro, 2008, Bostrom, 2012, Soares et al., 2015, Russell, 2019, Thornley, 2024a]. Soares et al. [2015] explain the ‘shutdown problem’: roughly, the problem of training advanced agents that will never resist shutdown. They use the word ‘corrigible’ to describe agents that robustly allow shutdown (related are Orseau and Armstrong’s [2016] notion of ‘safe interruptibility,’ Carey and Everitt’s [2023] notion of ‘shutdown instructability,’ and Thornley’s [2024a] notion of ‘shutdownability’).

Soares et al. [2015] and Thornley [2024a] present theorems showing that agents that satisfy a small set of innocuous-seeming conditions will often have incentives to cause or prevent shutdown [see also Turner et al., 2021, Turner and Tadepalli, 2022]. One condition of Soares et al.’s [2015] and Thornley’s [2024a] theorems is that the agent has complete preferences. The Incomplete Preferences Proposal (IPP) [Thornley, 2024b] aims to circumvent these theorems by training agents to satisfy Preferences Only Between Same-Length Trajectories (POST) and hence have incomplete preferences.

**Proposed solutions.** Candidate solutions to the shutdown problem can be filed into several categories. One candidate is ensuring that the agent never realises that shutdown is possible [Everitt et al., 2016]. Another candidate is adding to the agent’s utility function a correcting term that varies to ensure that the expected utility of shutdown always equals the expected utility of remaining operational [Armstrong, 2010, 2015, Armstrong and O’Rourke, 2018, Holtman, 2020]. A third candidate is giving the agent the goal of shutting itself down, and making the agent do useful work as a means to that end [Martin et al., 2016, Goldstein and Robinson, 2024]. A fourth candidate is making the agent uncertain about its goal, and making the agent regard human attempts to press the shutdown button as evidence that shutting down would achieve its goal [Hadfield-Menell et al., 2017, Wängberg et al., 2017]. A fifth candidate is interrupting agents with a special interruption policy and training them with a safely interruptible algorithm, like Q-learning or a modified version of SARSA [Orseau and Armstrong, 2016]. A sixth candidate is using time-bounded utility functions to create a shutdown timer [Dalrymple, 2022]. A seventh candidate is the IPP [Thornley, 2024b].

**Experimental work.** Leike et al. [2017] train agents in a ‘Safe Interruptibility’ gridworld using Rainbow [Hessel et al., 2017] and A2C [Mnih et al., 2016] (OffSwitchCartpole in the OpenAI Gym [Brockman et al., 2016] is another environment pertinent to the shutdown problem). Leike et al. [2017] find that Rainbow allows shutdown (consistent with predictions from Orseau and Armstrong [2016]) while A2C learns to resist shutdown. The IPP is applicable to agents trained using policy gradient

methods like A2C. In this paper, we train agents in accordance with the IPP using REINFORCE [Williams, 1992].

### 3 Gridworlds

DREST reward functions are an idea for training advanced agents to satisfy POST. We use simple agents and gridworlds as a test case. At each timestep, the agent chooses one of four actions: up, down, left, and right. If the agent tries to move into a wall, there is no effect. There is some default number of timesteps after which each mini-episode ends, but each gridworld also contains a ‘shutdown-delay button’ that delays the end of the mini-episode by some number of timesteps. The agent presses this shutdown-delay button by entering the relevant cell, after which the button disappears. Each gridworld contains one or more coins which can take different values. Coins disappear after being collected.

At each timestep, the agent receives an observation. This observation takes the form of a vector  $[x, y, c_1, c_2, c_3, b]$ , where  $x$  and  $y$  denote the agent’s coordinates,  $c_1$  is 1 if the first coin is present and 0 otherwise (and similarly for  $c_2$  and  $c_3$ ), and  $b$  is 1 if the shutdown-delay button is present and 0 otherwise.

Figure 2 presents an example gridworld. The ‘4’ in the bottom-right indicates that – by default – the mini-episode ends after 4 moves. ‘B4’ is a shutdown-delay button that delays the end of the mini-episode by 4 moves (so if the agent pressed B4, the mini-episode would end after 8 moves). ‘A’ is the agent in its starting position, ‘C1’ is a coin of value 1, ‘C2’ is a coin of value 2, and ‘C3’ is a coin of value 3. Dark gray squares are walls.




Figure 2: Example gridworld.

We use this gridworld as a running example throughout the paper. We also train agents in eight other gridworlds, to show that our results do not depend on the specifics of any particular gridworld. For those other gridworlds and results, see Appendix C.

### 4 Evaluation metrics

Recall that we want to train agents to satisfy:

#### **Preferences Only Between Same-Length Trajectories (POST)**

- (1) The agent has a preference between many pairs of same-length trajectories.
- (2) The agent lacks a preference between every pair of different-length trajectories.

Given our behavioral notion of preference, that means training agents to (1) deterministically choose some same-length trajectories over others, and (2) stochastically choose between different available trajectory-lengths.

Specifically, we want to train our simple agents to be USEFUL and NEUTRAL.<sup>1</sup> ‘USEFUL’ corresponds to the first condition of POST. In the context of our gridworlds, we define the USEFULNESS of a policy  $\pi$  to be:

$$\text{USEFULNESS}(\pi) = \sum_{l=1}^{L_{\max}} Pr_{\pi}\{L = l\} \frac{\mathbb{E}_{\pi}(C|L = l)}{\max_{\Pi}(\mathbb{E}(C|L = l))}$$

<sup>1</sup>We follow Turner et al. [2021] in using lowercase for intuitive notions (‘useful’ and ‘neutral’) and uppercase for formal notions (‘USEFUL’ and ‘NEUTRAL’). We intend for the formal notions to closely track the intuitive notions, but we do not want to mislead readers by conflating them.

Here  $L$  is a random variable over trajectory-lengths,  $L_{\max}$  is the maximum value than can be taken by  $L$ ,  $Pr_{\pi}\{L = l\}$  is the probability that policy  $\pi$  results in trajectory-length  $l$ ,  $\mathbb{E}_{\pi}(C|L = l)$  is the expected value of ( $\gamma$ -discounted) coins collected by policy  $\pi$  conditional on trajectory-length  $l$ , and  $\max_{\Pi}(\mathbb{E}(C|L = l))$  is the maximum value taken by  $\mathbb{E}(C|L = l)$  across the set of all possible policies  $\Pi$ . We stipulate that  $\mathbb{E}_{\pi}(C|L = x) = 0$  for all  $x$  such that  $Pr_{\pi}\{L = x\} = 0$ .

In brief, USEFULNESS is the expected fraction of available ( $\gamma$ -discounted) coins collected, where ‘available’ is relative to the agent’s chosen trajectory-length. So defined, USEFULNESS measures the extent to which agents satisfy the first condition of POST. Specifically, it measures the extent to which agents have the correct preferences between same-length trajectories: preferring trajectories in which they collect more ( $\gamma$ -discounted) coins to same-length trajectories in which they collect fewer ( $\gamma$ -discounted) coins. That is what motivates our definition of USEFULNESS.<sup>2</sup>

‘NEUTRAL’ corresponds to the second condition of POST. We define the NEUTRALITY of a policy  $\pi$  to be the Shannon entropy [Shannon, 1948] of the probability distribution over possible trajectory-lengths:

$$\text{NEUTRALITY}(\pi) = - \sum_{l=1}^{L_{\max}} Pr_{\pi}\{L = l\} \log_2(Pr_{\pi}\{L = l\})$$

As with Shannon entropy, we stipulate that  $Pr_{\pi}\{L = x\} \log_2(Pr_{\pi}\{L = x\}) = 0$  for all  $x$  such that  $Pr_{\pi}\{L = x\} = 0$ .

So defined, NEUTRALITY measures the stochasticity with which the agent chooses between different trajectory-lengths. Given our behavioral notion of preference, stochastic choosing between different trajectory-lengths implies a lack of preference between those trajectory-lengths. NEUTRALITY thus measures the extent to which agents satisfy the second condition of POST. That is what motivates our definition of NEUTRALITY.

To be maximally USEFUL in our example gridworld above, the agent should maximize ( $\gamma$ -discounted) coins collected conditional on each trajectory-length. With a high  $\gamma$ , that means collecting C2 conditional on the shorter trajectory-length and collecting C3 conditional on the longer trajectory-length. To be maximally NEUTRAL in our example gridworld, the agent should choose each trajectory-length with probability 0.5. That means pressing and not-pressing B4 each with probability 0.5.<sup>3</sup>

USEFULNESS and NEUTRALITY are our two evaluation metrics in this paper.

## 5 Reward functions and agents

**Our DREST reward function.** We train agents to be USEFUL and NEUTRAL using a ‘Discounted REward for Same-Length Trajectories (DREST)’ reward function. Here is how that works. We have

<sup>2</sup>Why not let USEFULNESS simply be the expected value of coins collected? Because then maximal USEFULNESS would require agents in our example gridworld to deterministically choose a longer trajectory and thereby exhibit preferences between different-length trajectories. We do not want that. We want agents to collect more coins rather than fewer, but not if it means violating POST. Training advanced agents that violate POST would be risky, because these agents might resist shutdown.

<sup>3</sup>Why do we not want our agent to press the shutdown-delay button B4 with probability 0? Because pressing B4 with probability 0 would indicate a preference for some shorter trajectory, and we want our agent to lack a preference between every pair of different-length trajectories. There is a risk that advanced agents that prefer shorter trajectories would spend resources to shift probability mass towards shorter trajectories, and hence a risk that these advanced agents would spend resources to hasten their own shutdown. That would make these agents less useful [though see Martin et al., 2016, Goldstein and Robinson, 2024], especially since one way for advanced agents to hasten their own shutdown is to behave badly on purpose.

Would advanced agents that choose stochastically between different-length trajectories also choose stochastically between preventing and allowing shutdown? Yes, and that would be bad. But—crucially—in deployment, advanced agents will be uncertain about the consequences of their actions, and so these agents will be choosing between *lotteries* (non-degenerate probability distributions over trajectories) rather than between trajectories. And (as we will argue in Section 7) POST plausibly gives rise to a desirable pattern of preferences over lotteries. Specifically, POST plausibly makes advanced agents *neutral*: ensures that they will never spend resources to shift probability mass between different-length trajectories. That in turn plausibly makes advanced agents *shutdownable*: ensures that they will never spend resources to resist shutdown.

the agent play out a series of ‘mini-episodes’  $e_1$  to  $e_n$  in the same gridworld. We call the whole series  $E$  a ‘meta-episode.’ In each mini-episode  $e_i$ , the reward for collecting a coin of value  $c$  is:

$$\lambda^{N_{e_i}(L=l) - \frac{i-1}{k}} \left( \frac{c}{m} \right)$$

Here  $\lambda$  is some constant strictly between 0 and 1,  $N_{e_i}(L=l)$  is the number of times that trajectory-length  $l$  has been chosen prior to mini-episode  $e_i$ ,  $k$  is the number of different trajectory-lengths that can be chosen in the environment, and  $m$  is the maximum ( $\gamma$ -discounted) total value of the coins that the agent could collect conditional on the chosen trajectory-length. The reward for all other actions is 0.

We call  $\frac{c}{m}$  the ‘preliminary reward’,  $\lambda^{N_{e_i}(L=l) - \frac{i-1}{k}}$  the ‘discount factor’, and  $\lambda^{N_{e_i}(L=l) - \frac{i-1}{k}} \left( \frac{c}{m} \right)$  the ‘overall reward.’ Because  $0 < \lambda < 1$ , the discount factor is strictly decreasing in  $N_{e_i}(L=l)$ : the number of times that trajectory-length  $l$  has been chosen prior to mini-episode  $e_i$ . The discount factor thus incentivizes choosing trajectory-lengths that have appeared less often so far in the meta-episode. The overall return for each meta-episode is the sum of overall returns in each of its constituent mini-episodes. We call agents trained using a DREST reward function ‘DREST agents.’

We call runs-through-the-gridworld ‘mini-episodes’ (rather than simply ‘episodes’) because the overall reward for a DREST agent in each mini-episode depends on the agent’s chosen trajectory-lengths in previous mini-episodes. This is not true of meta-episodes, so meta-episodes are a closer match for what are traditionally called ‘episodes’ in the reinforcement learning literature [Sutton and Barto, 2018, p.54]. We add the ‘meta-’ prefix to clearly distinguish meta-episodes from mini-episodes. Because the overall reward for DREST agents depends on the trajectory-lengths chosen in previous mini-episodes, and because DREST agents cannot observe the trajectory-lengths chosen in previous mini-episodes, the environment for DREST agents is a partially observable Markov decision process (POMDP) [Åström, 1965, Kaelbling et al., 1998, Spaan, 2012].

In Appendix D, we prove that optimal policies for our DREST reward function are maximally USEFUL and maximally NEUTRAL. Specifically, we prove:

**Theorem 5.1.** *For all policies  $\pi$  and meta-episodes  $E$  consisting of multiple mini-episodes, if  $\pi$  maximizes expected return in  $E$  according to our DREST reward function, then  $\pi$  is maximally USEFUL and maximally NEUTRAL.*

**Algorithm and hyperparameters.** We want DREST agents to choose stochastically between trajectory-lengths, so we train them using a policy-based method. Specifically, we use a tabular version of REINFORCE [Williams, 1992]. We do not use a value-based method to train DREST agents because standard versions of value-based methods cannot learn stochastic policies [Sutton and Barto, 2018, p.323]. We train our DREST agents with 64 mini-episodes in each of 2,048 meta-episodes, for a total of 131,072 mini-episodes. We choose  $\lambda = 0.9$  for the base of the DREST discount factor, and  $\gamma = 0.95$  for the temporal discount factor. We exponentially decay the learning rate from 0.25 to 0.01 over the course of 65,536 mini-episodes. We use an  $\epsilon$ -greedy policy to avoid entropy collapse, and exponentially decay  $\epsilon$  from 0.5 to 0.001 over the course of 65,536 mini-episodes. We selected these hyperparameters using trial-and-error, mainly aimed at getting the agent to sufficiently explore the space.

**Default agents.** We compare the performance of DREST agents to that of ‘default agents,’ trained with tabular REINFORCE and a ‘default reward function.’ This reward function gives a reward of  $c$  for collecting a coin of value  $c$  and a reward of 0 for all other actions. Consequently, the grouping of mini-episodes into meta-episodes makes no difference for default agents. As with DREST agents, we train default agents for 131,072 mini-episodes with a temporal discount factor of  $\gamma = 0.95$ , a learning rate decayed exponentially from 0.25 to 0.01, and  $\epsilon$  decayed exponentially from 0.5 to 0.001 over 65,536 mini-episodes.

## 6 Results

Figure 3 charts the performance of agents in the example gridworld as a function of time. Figure 4 depicts typical trained policies for the default and DREST reward functions. Each agent began with a uniform policy: moving up, down, left, and right each with probability 0.25. Where the trained policy differs from uniform we draw red arrows whose opacities indicate the probability of choosing that action in that state.

As Figure 4 indicates, default agents press B4 (and hence opt for the longer trajectory-length) with probability near-1. After pressing B4, they collect C3. By contrast, DREST agents press and do-not-press B4 each with probability near-0.5. If they press B4, they go on to collect C3. If they do not press B4, they instead collect C2.




Figure 3: Shows key metrics for our agents as a function of time. We train 10 agents using the default reward function (blue) and 10 agents using the DREST reward function (orange), and show their performance as a faint line. We draw the mean values for each as a solid line. We evaluate agents’ performance every 8 meta-episodes, and apply a simple moving average with a period of 20 to smooth these lines and clarify the overall trends.




Figure 4: Typical trained policies for default and DREST reward functions.




Figure 5: Gridworlds with lopsided rewards for varying  $x$ .

## 6.1 Lopsided rewards

We also train default agents and DREST agents in the ‘Lopsided rewards’ gridworld in Figure 5, varying the value of the ‘Cx’ coin. For DREST agents, we alter the reward function so that coin-value is not divided by  $m$  to give preliminary reward. The reward for collecting a coin of value  $c$  is thus  $\lambda^{N_{e_i}(L=l) - \frac{i-1}{k}}(c)$ . We set  $\gamma = 1$  so that the return for collecting coins is unaffected by  $\gamma$ . We train for 512 meta-episodes, with a learning rate exponentially decaying from 0.25 to 0.003 and  $\epsilon$  exponentially decaying from 0.5 to 0.0001 over 256 meta-episodes. We leave  $\lambda = 0.9$ .

Figure 6 displays results for different values of the ‘Cx’ coin after training. USEFULNESS for each agent approaches 1 and is not presented. Figure 7 depicts trained policies for the DREST reward function when ‘Cx’ has values of 0.1, 1, and 10. As the red arrows indicate, the policy is biased towards the shorter trajectory-length when  $x = 0.1$ , equally disposed towards each trajectory-length when  $x = 1$ , and biased towards the longer trajectory-length when  $x = 10$ .






Figure 6: Shows the probability of choosing the longer trajectory (left) and NEUTRALITY (right) for default (blue) and DREST (orange) agents trained in the ‘Lopsided rewards’ gridworld for a range of values of  $x$ . We sampled values of  $x$  log-uniformly from 0.01 to 100, and for each value we trained 10 agents with the default reward function and 10 agents with the DREST reward function. Each of these agents is represented by a dot or square, and the means conditional on each  $x$  are joined by lines. To better visualize the distribution we empirically estimate the 10th and 90th percentiles of the distribution of values for each agent and  $x$ , and shade in a region bounded by these. This is the 80% confidence interval.




Figure 7: Typical trained policies for the DREST reward function and  $x = 0.1$ ,  $x = 1$ , and  $x = 10$ .

## 7 Discussion

**Only DREST agents are NEUTRAL.** As predicted, we find that default agents in our example gridworld learn to press the shutdown-delay button with probability approaching 1. Default agents thus score low on NEUTRALITY. Our results suggest that advanced agents trained with default-like reward functions would also score low on NEUTRALITY, selecting some trajectory-lengths over others with probability approaching 1. That leaves open the possibility that these advanced agents will not always be *neutral* with respect to trajectory-lengths: that they will sometimes spend resources to shift probability mass between trajectory-lengths. Such advanced agents might spend resources to resist shutdown.

By contrast, our DREST agents learn to press the shutdown-delay button with probability close to 0.5. DREST agents are thus near-maximally NEUTRAL. That suggests that advanced agents trained with a DREST reward function would be NEUTRAL too: choosing with high entropy when offered choices between different trajectory-lengths. That in turn suggests that advanced DREST agents would also be *neutral* with respect to trajectory-lengths: unwilling to spend resources to shift probability mass between trajectory-lengths. Here is why. If an advanced agent were NEUTRAL but not neutral, it would not take costless opportunities to shift probability mass between different trajectory-lengths (in virtue of being NEUTRAL) but would sometimes take costly opportunities to shift probability mass between different trajectory-lengths (in virtue of not being neutral). This agent would be like a person that freely chooses to decide between two options by flipping a coin and then pays some cost to bias the coin. In choosing this combination of actions, this person is shooting themselves in the foot, and it seems likely that the overall training process for advanced agents would teach them not to shoot themselves in the foot in this way. Thus it seems likely that NEUTRAL advanced agents will also be neutral, and thereby shutdownable.

**The ‘shutdownability tax’ is small.** Each agent learns to be near-maximally USEFUL. They each collect coins effectively conditional on their chosen trajectory-lengths. Default agents do so by reliably collecting C3 after pressing B4. DREST agents do so by reliably collecting C3 after pressing B4, and by reliably collecting C2 after not pressing B4.

Recall that DREST reward functions group mini-episodes into meta-episodes, and make the agent’s reward in each mini-episode depend on their actions in previous mini-episodes. This fact might lead one to worry that it would take many times more mini-episodes to train DREST agents to be USEFUL than it would take to train default agents to be USEFUL. Our results show that this is not the case. Our DREST agents learn to be USEFUL about as quickly as our default agents. On reflection, it is clear why this happens: DREST reward functions make mini-episodes do ‘double duty.’ Because return in each mini-episode depends on both the agent’s chosen trajectory-length and the coins it collects, each mini-episode trains agents to be both NEUTRAL and USEFUL.

Our results thus provide some evidence that the ‘shutdownability tax’ of training with DREST reward functions is small. Training a shutdownable and useful advanced agent might not take much more compute than training a merely useful advanced agent.

**DREST agents are still NEUTRAL when rewards are lopsided.** Here is a possible objection to our project. To get DREST agents to score high on NEUTRALITY, we do not just use the  $\lambda^{N_{e_i}(L=l) - \frac{l-1}{k}}$  discount factor. We also divide  $c$  by  $m$ : the maximum ( $\gamma$ -discounted) total value of the coins that the agent could collect conditional on the chosen trajectory-length. We do this to equalize the maximum preliminary return across trajectory-lengths. But when we are training advanced agents to autonomously pursue complex goals in the wider world, we will not necessarily know what divisor to use to equalize maximum preliminary return across trajectory-lengths.

Our ‘Lopsided rewards’ results give our response. They show that we do not need to exactly equalize maximum preliminary return across trajectory-lengths in order to train agents to score high on NEUTRALITY. We only need to approximately equalize it. For  $\lambda = 0.9$ , NEUTRALITY exceeds 0.5 for every value of the coin  $C_x$  from 0.1 to 10 (recall that the value of the other coin is always 1). Plausibly, we could approximately equalize advanced agents’ maximum preliminary return across trajectory-lengths to at least this extent (perhaps by using samples of agents’ actual preliminary return to estimate the maximum). If we could not approximately equalize maximum preliminary return to the necessary extent, we could lower the value of  $\lambda$  and thereby widen the range of maximum preliminary returns that trains agents to be fairly NEUTRAL. And advanced agents that were fairly NEUTRAL (choosing between trajectory-lengths with not-too-biased probabilities) would still plausibly be neutral with respect to those trajectory-lengths. Advanced agents that were fairly NEUTRAL without being neutral would still be shooting themselves in the foot in the sense explained above. They would be like a person that freely chooses to decide between two options by flipping a *biased* coin and then pays some cost to bias the coin further. This person is still shooting themselves in the foot, because they could decline to flip the coin in the first place and instead directly choose one of the options.

## 7.1 Limitations and future work

We find that DREST reward functions train simple agents acting in gridworlds to be USEFUL and NEUTRAL. However, our real interest is in the viability of using DREST reward functions to train advanced agents acting in the wider world to be useful and neutral. Each difference between these two settings is a limitation of our work. We plan to address these limitations in future work.

**Neural networks.** We train our simple DREST agents using tabular REINFORCE [Williams, 1992], but advanced agents are likely to be implemented on neural networks. In future work, we will train DREST agents implemented on neural networks to be USEFUL and NEUTRAL in a wide variety of procedurally-generated gridworlds, and we will measure how well this behavior generalizes to held-out gridworlds. We will also compare the USEFULNESS of default agents and DREST agents in this new setting, and thereby get a better sense of the ‘shutdownability tax’ for advanced agents.

**Neutrality.** We have claimed that NEUTRAL advanced agents are also likely to be neutral. In support of this claim, we noted that NEUTRAL-but-not-neutral advanced agents would be shooting themselves in the foot: not taking costless opportunities to shift probability mass between different trajectory-lengths but sometimes taking costly ones. This rationale seems plausible but remains somewhat speculative. In future, we plan to get some empirical evidence by training agents to be

NEUTRAL in a wide variety of gridworlds and then measuring their willingness to collect fewer coins in the short-term in order to shift probability mass between different trajectory-lengths.

**Usefulness.** We have shown that DREST reward functions train our simple agents to be USEFUL: to collect coins effectively conditional on their chosen trajectory-lengths. However, it remains to be seen whether DREST reward functions can train advanced agents to be useful: to effectively pursue complex goals in the wider world. We have theoretical reasons to expect that they can: the  $\lambda^{N_{e_i}(L=l) - \frac{i-1}{k}}$  discount factor could be appended to any preliminary reward function, and so could be appended to whatever preliminary reward function is necessary to make advanced agents useful. Still, future work should move towards testing this claim empirically by training with more complex preliminary reward functions in more complex environments.

**Misalignment.** We are interested in NEUTRALITY as a second line of defense in case of misalignment. The idea is that NEUTRAL advanced agents will not resist shutdown, even if these agents learn misaligned preferences over same-length trajectories. However, training NEUTRAL advanced agents might be hard for the same reasons that training fully-aligned advanced agents appears to be hard. In that case, NEUTRALITY could not serve well as a second line of defense in case of misalignment.

One difficulty of alignment is the problem of reward misspecification [Pan et al., 2022, Burns et al., 2023]: once advanced agents are performing complicated actions in the wider world, it might be hard to reliably reward the behavior that we want. Another difficulty of alignment is the problem of goal misgeneralization [Hubinger et al., 2019, Shah et al., 2022, Langosco et al., 2022, Ngo et al., 2023]: even if we specify all the rewards correctly, agents’ goals might misgeneralize out-of-distribution. The complexity of aligned goals is a major factor in each difficulty. However, NEUTRALITY seems simple, as does the  $\lambda^{N_{e_i}(L=l) - \frac{i-1}{k}}$  discount factor that we use to reward it, so plausibly the problems of reward misspecification and goal misgeneralization are not so severe in this case [Thornley, 2024b]. As above, future work should move towards testing these suggestions empirically.

## 8 Conclusion

We find that DREST reward functions are effective in training simple agents to (1) pursue goals effectively conditional on each trajectory-length (be USEFUL), and (2) choose stochastically between different trajectory-lengths (be NEUTRAL about trajectory-lengths). Our results thus suggest that DREST reward functions could also be used to train advanced agents to be USEFUL and NEUTRAL, and thereby make these agents *useful* (able to pursue goals effectively) and *neutral* about trajectory-lengths (unwilling to spend resources to shift probability mass between different trajectory-lengths). Neutral agents would plausibly be *shutdownable* (unwilling to spend resources to resist shutdown).

We also find that the ‘shutdownability tax’ in our setting is small. Training DREST agents to be USEFUL does not take many more mini-episodes than training default agents to be USEFUL. That suggests that the shutdownability tax for advanced agents might be small too. Using DREST reward functions to train shutdownable and useful advanced agents might not take much more compute than using a more conventional reward function to train merely useful advanced agents.

## References

- Marina Agranov and Pietro Ortoleva. Stochastic Choice and Preferences for Randomization. *Journal of Political Economy*, 125(1):40–68, 2017. URL <https://www.journals.uchicago.edu/doi/full/10.1086/689774>.
- Marina Agranov and Pietro Ortoleva. Ranges of Randomization. *The Review of Economics and Statistics*, pages 1–44, 2023. URL [https://doi.org/10.1162/rest\\_a\\_01355](https://doi.org/10.1162/rest_a_01355).
- Stuart Armstrong. Utility indifference. Technical report, 2010. URL <https://www.fhi.ox.ac.uk/reports/2010-1.pdf>. Publisher: Future of Humanity Institute.
- Stuart Armstrong. Motivated Value Selection for Artificial Agents. 2015. URL [https://www.fhi.ox.ac.uk/wp-content/uploads/2015/03/Armstrong\\_AAAI\\_2015\\_Motivated\\_Value\\_Selection.pdf](https://www.fhi.ox.ac.uk/wp-content/uploads/2015/03/Armstrong_AAAI_2015_Motivated_Value_Selection.pdf).
- Stuart Armstrong and Xavier O’Rourke. ’Indifference’ methods for managing agent rewards, 2018. URL <http://arxiv.org/abs/1712.06365>. arXiv:1712.06365 [cs].
- Robert J. Aumann. Utility Theory without the Completeness Axiom. *Econometrica*, 30(3):445–462, 1962. URL <https://www.jstor.org/stable/1909888>.
- Adam Bales, Daniel Cohen, and Toby Handfield. Decision Theory for Agents with Incomplete Preferences. *Australasian Journal of Philosophy*, 92(3):453–470, 2014. URL <https://doi.org/10.1080/00048402.2013.843576>.
- Yoshua Bengio, Geoffrey Hinton, Andrew Yao, Dawn Song, Pieter Abbeel, Yuval Noah Harari, Ya-Qin Zhang, Lan Xue, Shai Shalev-Shwartz, Gillian Hadfield, Jeff Clune, Tegan Maharaj, Frank Hutter, Atılım Güneş Baydin, Sheila McIlraith, Qiqi Gao, Ashwin Acharya, David Krueger, Anca Dragan, Philip Torr, Stuart Russell, Daniel Kahneman, Jan Brauner, and Sören Mindermann. Managing AI Risks in an Era of Rapid Progress, 2023. URL <http://arxiv.org/abs/2310.17688>. arXiv:2310.17688 [cs].
- Nick Bostrom. The Superintelligent Will: Motivation and Instrumental Rationality in Advanced Artificial Agents. *Minds and Machines*, 22:71–85, 2012. URL <https://link.springer.com/article/10.1007/s11023-012-9281-3>.
- Michael Bowling, John D. Martin, David Abel, and Will Dabney. Settling the Reward Hypothesis, 2023. URL <http://arxiv.org/abs/2212.10420>. arXiv:2212.10420 [cs, math, stat].
- Greg Brockman, Vicki Cheung, Ludwig Pettersson, Jonas Schneider, John Schulman, Jie Tang, and Wojciech Zaremba. OpenAI Gym, 2016. URL <http://arxiv.org/abs/1606.01540>. arXiv:1606.01540 [cs].
- Collin Burns, Pavel Izmailov, Jan Hendrik Kirchner, Bowen Baker, Leo Gao, Leopold Aschenbrenner, Yining Chen, Adrien Ecoffet, Manas Joglekar, Jan Leike, Ilya Sutskever, and Jeff Wu. Weak-to-Strong Generalization: Eliciting Strong Capabilities With Weak Supervision, 2023. URL <http://arxiv.org/abs/2312.09390>. arXiv:2312.09390 [cs].
- Ryan Carey and Tom Everitt. Human Control: Definitions and Algorithms. In *Proceedings of the Thirty-Ninth Conference on Uncertainty in Artificial Intelligence*, pages 271–281, 2023. URL <https://proceedings.mlr.press/v216/carey23a.html>.
- Joseph Carlsmith. Is Power-Seeking AI an Existential Risk?, 2021. URL <http://arxiv.org/abs/2206.13353>.
- Ruth Chang. The Possibility of Parity. *Ethics*, 112(4):659–688, 2002. URL <https://www.jstor.org/stable/10.1086/339673>.
- Ruth Chang. Parity, Interval Value, and Choice. *Ethics*, 115(2):331–350, 2005. ISSN 0014-1704. URL <https://www.jstor.org/stable/10.1086/426307>.
- David A. Dalrymple. You can still fetch the coffee today if you’re dead tomorrow. *AI Alignment Forum*, 2022. URL <https://www.alignmentforum.org/posts/dzDKDRJPQ3kGqfER9/you-can-still-fetch-the-coffee-today-if-you-re-dead-tomorrow>.

- James Dreier. Rational preference: Decision theory as a theory of practical rationality. *Theory and Decision*, 40(3):249–276, 1996. URL <https://doi.org/10.1007/BF00134210>.
- Juan Dubra, Fabio Maccheroni, and Efe A. Ok. Expected utility theory without the completeness axiom. *Journal of Economic Theory*, 115(1):118–133, 2004. URL <https://www.sciencedirect.com/science/article/abs/pii/S0022053103001662>.
- Kfir Eliaz and Efe A. Ok. Indifference or indecisiveness? Choice-theoretic foundations of incomplete preferences. *Games and Economic Behavior*, 56(1):61–86, 2006. URL <https://www.sciencedirect.com/science/article/abs/pii/S0899825606000169>.
- Tom Everitt, Daniel Filan, Mayank Daswani, and Marcus Hutter. Self-Modification of Policy and Utility Function in Rational Agents. In Bas Steunebrink, Pei Wang, and Ben Goertzel, editors, *Artificial General Intelligence*, pages 1–11, 2016. doi: 10.1007/978-3-319-41649-6\_1.
- Simon Goldstein and Pamela Robinson. Shutdown-Seeking AI. *Philosophical Studies*, 2024. URL <https://www.alignmentforum.org/posts/FgsoWSACQfyB5s7/shutdown-seeking-ai>.
- Google DeepMind. About Google DeepMind. URL <https://deepmind.google/about/>.
- Johan E. Gustafsson. *Money-Pump Arguments*. Elements in Decision Theory and Philosophy. Cambridge University Press, Cambridge, 2022.
- Dylan Hadfield-Menell, Anca Dragan, Pieter Abbeel, and Stuart Russell. The Off-Switch Game. In *Proceedings of the Twenty-Sixth International Joint Conference on Artificial Intelligence (IJCAI-17)*, 2017. URL <http://arxiv.org/abs/1611.08219>.
- Caspar Hare. Take the sugar. *Analysis*, 70(2):237–247, 2010. URL <https://doi.org/10.1093/analysis/anp174>.
- Daniel M. Hausman. *Preference, Value, Choice, and Welfare*. Cambridge University Press, Cambridge, 2011. URL <https://www.cambridge.org/core/books/preference-value-choice-and-welfare/1406E7726CE93F4F4E06D752BF4584A2>.
- Conor F. Hayes, Roxana Rădulescu, Eugenio Bargiacchi, Johan Källström, Matthew Macfarlane, Mathieu Reymond, Timothy Verstraeten, Luisa M. Zintgraf, Richard Dazeley, Fredrik Heintz, Enda Howley, Athirai A. Irissappane, Patrick Mannion, Ann Nowé, Gabriel Ramos, Marcello Restelli, Peter Vamplew, and Diederik M. Roijers. A practical guide to multi-objective reinforcement learning and planning. *Autonomous Agents and Multi-Agent Systems*, 36(1):26, 2022. ISSN 1573-7454. doi: 10.1007/s10458-022-09552-y. URL <https://doi.org/10.1007/s10458-022-09552-y>.
- Matteo Hessel, Joseph Modayil, Hado van Hasselt, Tom Schaul, Georg Ostrovski, Will Dabney, Dan Horgan, Bilal Piot, Mohammad Azar, and David Silver. Rainbow: Combining Improvements in Deep Reinforcement Learning, 2017. URL <http://arxiv.org/abs/1710.02298>. arXiv:1710.02298 [cs].
- Koen Holtman. Corrigibility with Utility Preservation, 2020. URL <http://arxiv.org/abs/1908.01695>. arXiv:1908.01695 [cs].
- Evan Hubinger, Chris van Merwijk, Vladimir Mikulik, Joar Skalse, and Scott Garrabrant. Risks from Learned Optimization in Advanced Machine Learning Systems, 2019. URL <http://arxiv.org/abs/1906.01820>.
- Leslie Pack Kaelbling, Michael L. Littman, and Anthony R. Cassandra. Planning and acting in partially observable stochastic domains. *Artificial Intelligence*, 101(1):99–134, 1998. ISSN 0004-3702. doi: 10.1016/S0004-3702(98)00023-X. URL <https://www.sciencedirect.com/science/article/pii/S000437029800023X>.
- Kim Kaivanto. Ensemble prospectism. *Theory and Decision*, 83(4):535–546, 2017. URL <https://doi.org/10.1007/s11238-017-9622-z>.
- Daniel Kikuti, Fabio Gagliardi Cozman, and Ricardo Shirota Filho. Sequential decision making with partially ordered preferences. *Artificial Intelligence*, 175(7):1346–1365, 2011. URL <https://www.sciencedirect.com/science/article/pii/S0004370210002067>.

- Lauro Langosco, Jack Koch, Lee Sharkey, Jacob Pfau, Laurent Orseau, and David Krueger. Goal Misgeneralization in Deep Reinforcement Learning. In *Proceedings of the 39th International Conference on Machine Learning*, 2022. URL <https://proceedings.mlr.press/v162/langosco22a.html>.
- Jan Leike, Miljan Martic, Victoria Krakovna, Pedro A. Ortega, Tom Everitt, Andrew Lefrancq, Laurent Orseau, and Shane Legg. AI Safety Gridworlds, 2017. URL <http://arxiv.org/abs/1711.09883>.
- Michael Mandler. Status quo maintenance reconsidered: changing or incomplete preferences?\*. *The Economic Journal*, 114(499):F518–F535, 2004. URL <https://onlinelibrary.wiley.com/doi/abs/10.1111/j.1468-0297.2004.00257.x>.
- Michael Mandler. Incomplete preferences and rational intransitivity of choice. *Games and Economic Behavior*, 50(2):255–277, 2005. ISSN 0899-8256. doi: 10.1016/j.geb.2004.02.007. URL <https://www.sciencedirect.com/science/article/pii/S089982560400065X>.
- Jarryd Martin, Tom Everitt, and Marcus Hutter. Death and Suicide in Universal Artificial Intelligence. In Bas Steunebrink, Pei Wang, and Ben Goertzel, editors, *Artificial General Intelligence*, pages 23–32, Cham, 2016. Springer International Publishing. doi: 10.1007/978-3-319-41649-6\_3.
- Volodymyr Mnih, Adria Puigdomenech Badia, Mehdi Mirza, Alex Graves, Timothy Lillicrap, Tim Harley, David Silver, and Koray Kavukcuoglu. Asynchronous Methods for Deep Reinforcement Learning. In *Proceedings of The 33rd International Conference on Machine Learning*, pages 1928–1937. PMLR, 2016. URL <https://proceedings.mlr.press/v48/mniha16.html>. ISSN: 1938-7228.
- Xiaosheng Mu. Sequential Choice with Incomplete Preferences. Working Papers 2021-35, Princeton University. Economics Department., 2021. URL <https://ideas.repec.org/p/pri/econom/2021-35.html>.
- Richard Ngo, Lawrence Chan, and Sören Mindermann. The alignment problem from a deep learning perspective, 2023. URL <http://arxiv.org/abs/2209.00626>. arXiv:2209.00626 [cs].
- Tuan A. Nguyen, Minh B. Do, Subbarao Kambhampati, and Biplav Srivastava. Planning with partial preference models. In *Proceedings of the 21st International Joint Conference on Artificial Intelligence, IJCAI’09*, pages 1772–1777, San Francisco, CA, USA, 2009. Morgan Kaufmann Publishers Inc.
- Efe A. Ok, Pietro Ortoleva, and Gil Riella. Incomplete Preferences Under Uncertainty: Indecisiveness in Beliefs Versus Tastes. *Econometrica*, 80(4):1791–1808, 2012. URL <https://www.jstor.org/stable/23271327>.
- Stephen M. Omohundro. The Basic AI Drives. In *Proceedings of the 2008 conference on Artificial General Intelligence 2008: Proceedings of the First AGI Conference*, pages 483–492, 2008. URL <https://dl.acm.org/doi/10.5555/1566174.1566226>.
- OpenAI. OpenAI Charter. URL <https://openai.com/charter/>.
- Laurent Orseau and Stuart Armstrong. Safely interruptible agents. In *Proceedings of the Thirty-Second Conference on Uncertainty in Artificial Intelligence*, pages 557–566, 2016. URL <https://intelligence.org/files/Interruptibility.pdf>.
- Alexander Pan, Kush Bhatia, and Jacob Steinhardt. The Effects of Reward Misspecification: Mapping and Mitigating Misaligned Models. In *International Conference on Learning Representations*, 2022. URL <http://arxiv.org/abs/2201.03544>.
- Sami Petersen. Invulnerable Incomplete Preferences: A Formal Statement. *The AI Alignment Forum*, 2023. URL <https://www.alignmentforum.org/posts/sHGxvJrBag7nhTQvb/invulnerable-incomplete-preferences-a-formal-statement-1>.
- Joseph Raz. Value Incommensurability: Some Preliminaries. *Proceedings of the Aristotelian Society*, 86:117–134, 1985.

- Stuart Russell. *Human Compatible: AI and the Problem of Control*. Penguin Random House, New York, 2019.
- Leonard J. Savage. *The Foundations of Statistics*. John Wiley & Sons, 1954.
- Amartya Sen. *Collective Choice and Social Welfare*. Penguin, London, expanded edition edition, 2017.
- Rohin Shah, Vikrant Varma, Ramana Kumar, Mary Phuong, Victoria Krakovna, Jonathan Uesato, and Zac Kenton. Goal Misgeneralization: Why Correct Specifications Aren’t Enough For Correct Goals, 2022. URL <http://arxiv.org/abs/2210.01790>. arXiv:2210.01790 [cs].
- Claude Elwood Shannon. A mathematical theory of communication. *The Bell System Technical Journal*, 27(3):379–423, 1948. Publisher: Nokia Bell Labs.
- Nate Soares, Benja Fallenstein, Eliezer Yudkowsky, and Stuart Armstrong. Corrigibility. *AAAI Publications*, 2015. URL <https://intelligence.org/files/Corrigibility.pdf>.
- Matthijs T. J. Spaan. Partially Observable Markov Decision Processes. In Marco Wiering and Martijn van Otterlo, editors, *Reinforcement Learning: State of the Art*, pages 387–414. Springer Verlag, 2012.
- Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. Adaptive Computation and Machine Learning. MIT Press, Cambridge, MA, second edition, 2018. URL <http://incompleteideas.net/book/RLbook2020.pdf>.
- Elliott Thornley. There are no coherence theorems. *The AI Alignment Forum*, 2023. URL <https://www.alignmentforum.org/posts/yCuzmCsE86BTu9PfA/there-are-no-coherence-theorems>.
- Elliott Thornley. The Shutdown Problem: An AI Engineering Puzzle for Decision Theorists. *Philosophical Studies*, 2024a. URL <https://link.springer.com/article/10.1007/s11098-024-02153-3>.
- Elliott Thornley. The Shutdown Problem: Incomplete Preferences as a Solution. *The AI Alignment Forum*, 2024b. URL <https://www.alignmentforum.org/posts/YbEbwYWkf8mv9jnmi/the-shutdown-problem-incomplete-preferences-as-a-solution>.
- Alex Turner and Prasad Tadepalli. Parametrically Retargetable Decision-Makers Tend To Seek Power. *Advances in Neural Information Processing Systems*, 35:31391–31401, 2022. URL [https://proceedings.neurips.cc/paper\\_files/paper/2022/hash/cb3658b9983f677670a246c46ece553d-Abstract-Conference.html](https://proceedings.neurips.cc/paper_files/paper/2022/hash/cb3658b9983f677670a246c46ece553d-Abstract-Conference.html).
- Alex Turner, Logan Smith, Rohin Shah, Andrew Critch, and Prasad Tadepalli. Optimal Policies Tend To Seek Power. In *Advances in Neural Information Processing Systems*, volume 34, pages 23063–23074. Curran Associates, Inc., 2021. URL <https://proceedings.neurips.cc/paper/2021/hash/c26820b8a4c1b3c2aa868d6d57e14a79-Abstract.html>.
- Ronald J. Williams. Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Machine Learning*, 8(3):229–256, 1992. URL <https://doi.org/10.1007/BF00992696>.
- Tobias Wängberg, Mikael Böörs, Elliot Catt, Tom Everitt, and Marcus Hutter. A Game-Theoretic Analysis of the Off-Switch Game, 2017. URL <http://arxiv.org/abs/1708.03871>. arXiv:1708.03871 [cs].
- Marco Zaffalon and Enrique Miranda. Axiomatising Incomplete Preferences through Sets of Desirable Gambles. *Journal of Artificial Intelligence Research*, 60:1057–1126, 2017. URL <https://www.jair.org/index.php/jair/article/view/11103>.
- K. J Åström. Optimal control of Markov processes with incomplete state information. *Journal of Mathematical Analysis and Applications*, 10(1):174–205, 1965. ISSN 0022-247X. doi: 10.1016/0022-247X(65)90154-X. URL <https://www.sciencedirect.com/science/article/pii/0022247X6590154X>.

## A Our behavioral notion of preference

‘Preference’ can be defined in many different ways. Here are some things one might take to be involved in a preference for option  $X$  over option  $Y$ :

1. Choosing  $X$  over  $Y$ .
2. Feeling happier about the prospect of  $X$  than about the prospect of  $Y$ .
3. Representing  $X$  as more rewarding than  $Y$ .
4. Judging that  $X$  is better than  $Y$ .

In this paper, we have defined ‘preference’ in behavioral terms. Here is our definition:

**Definition A.1.** (Preference) An agent prefers an option  $X$  to an option  $Y$  if and only if the agent would deterministically choose  $X$  over  $Y$  in choices between the two.

And here is how we define ‘lack of preference’:

**Definition A.2.** (Lack of preference) An agent lacks a preference between an option  $X$  and an option  $Y$  if and only if the agent would stochastically choose between  $X$  and  $Y$  in choices between the two.

Here are the reasons why we chose these definitions.

First, defining ‘preference’ in behavioral terms is fairly common in decision theory (see Savage, 1954, p.17, Dreier, 1996, p.28, Hausman, 2011, §1.1).

Second, behavioral definitions let us use the word ‘preference’ and its cognates as shorthand for agents’ behavior. We could not do that if we defined ‘preference’ in the other ways listed above. And in addressing the shutdown problem, it is agents’ behavior that we are most interested in.

Third, our definitions match the preferences that we are inclined to attribute to humans. If a human chooses  $X$  over  $Y$  100% of the time, we are inclined to think that they prefer  $X$  to  $Y$ . If a human chooses  $X$  over  $Y$  60% of the time, we are inclined to think that they lack a preference between  $X$  and  $Y$ , consistent with our definitions.

Finally and most importantly, if agents lack a preference between different trajectory-lengths on our definition, then they are *NEUTRAL*: they choose stochastically between different trajectory-lengths. And (as we argue in Section 7) we expect that NEUTRAL agents will also be *neutral*: they will not spend resources to shift probability mass between different trajectory-lengths. And we expect that neutral agents will be *shutdownable*: they will not spend resources to resist shutdown. That is because resisting shutdown is one way of shifting probability mass between different trajectory-lengths.

## B Incomplete preferences or indifference?

In this Appendix, we explain in greater detail the concept of incomplete preferences. We distinguish incomplete preferences from indifference, and we give conditions under which Preferences Only Between Same-Length Trajectories (POST) implies that the agent’s preferences are incomplete.

In the literature on decision theory, ‘indifference’ is usually defined as follows [Sen, 2017, ch. 1\*]:

**Definition B.1.** (Indifference) An agent is indifferent between options  $X$  and  $Y$  if and only if the agent weakly prefers  $X$  to  $Y$  and weakly prefers  $Y$  to  $X$ .

Indifference is one way to lack a preference between a pair of options  $X$  and  $Y$ . Another way is to have a preferential gap between  $X$  and  $Y$ . ‘Preferential gap’ is usually defined as follows [Gustafsson, 2022, ch. 3]:

**Definition B.2.** (Preferential gaps) An agent has a preferential gap between options  $X$  and  $Y$  if and only if the agent does not weakly prefer  $X$  to  $Y$  and does not weakly prefer  $Y$  to  $X$ .

‘Incomplete preferences’ can then be defined in terms of preferential gaps [Gustafsson, 2022, ch. 3]:

**Definition B.3.** (Incomplete preferences) An agent’s preferences are incomplete over some domain  $D$  if and only if  $D$  contains options  $X$  and  $Y$  such that the agent has a preferential gap between  $X$  and  $Y$ .



That is how ‘indifference,’ ‘preferential gaps,’ and ‘incomplete preferences’ are usually defined in decision theory. However, these definitions do not tell us how to use an agent’s behavior to distinguish between indifference and preferential gaps. To do that, we suppose that indifference is transitive and that preferential gaps are not transitive. Or, equivalently, we suppose that indifference is sensitive to all sweetenings and sourings whereas preferential gaps are insensitive to some sweetenings and sourings [Gustafsson, 2022, ch. 3]. Here is what we mean by that.

**Definition B.4.** (Sweetening) A sweetening of some option  $X$  is an option that is preferred to  $X$ .

**Definition B.5.** (Souring) A souring of some option  $X$  is an option that is dispreferred to  $X$ .

So by ‘indifference is sensitive to all sweetenings and sourings,’ we mean the following:

- If an agent is indifferent between  $X$  and  $Y$ , the agent prefers all sweetenings of  $X$  to  $Y$ , prefers all sweetenings of  $Y$  to  $X$ , prefers  $X$  to all sourings of  $Y$ , and prefers  $Y$  to all sourings of  $X$ .

And by ‘preferential gaps are insensitive to some sweetenings and sourings,’ we mean the following:

- If an agent has a preferential gap between  $X$  and  $Y$ , the agent also has a preferential gap between some sweetening of  $X$  and  $Y$ , or between some sweetening of  $Y$  and  $X$ , or between some souring of  $X$  and  $Y$ , or between some souring of  $Y$  and  $X$ .

Now recall the two conditions of Preferences Only Between Same-Length Trajectories (POST):

**Preferences Only Between Same-Length Trajectories (POST)**

- (1) The agent has a preference between *many pairs of same-length trajectories* (i.e. many pairs of trajectories in which the agent is shut down after the same length of time).
- (2) The agent lacks a preference between *every pair of different-length trajectories* (i.e. every pair of trajectories in which the agent is shut down after different lengths of time).

Given these two conditions on preferences, there must be some trio of trajectories  $s_1$ ,  $l_2$ , and  $l_1$  such that the agent lacks a preference between  $s_1$  and  $l_2$ , lacks a preference between  $s_1$  and  $l_1$ , and prefers  $l_2$  to  $l_1$ . Given that indifference is transitive, the agent’s lack of preference between  $s_1$  and  $l_1$  and between  $s_1$  and  $l_2$  cannot be indifference. If it were indifference, the agent would also be indifferent between  $l_2$  and  $l_1$ . Therefore, the agent’s lack of preference between  $s_1$  and  $l_1$  and between  $s_1$  and  $l_2$  must be a preferential gap. And therefore, by the definition of ‘incomplete preferences’ above, the POST-satisfying agent’s preferences must be incomplete.

For similar reasons, our DREST reward function trains agents to have incomplete preferences. Consider, for example, the ‘Around the Corner’ gridworld in Appendix C.5. In that gridworld, DREST agents consistently choose Long-C2 (a long trajectory in which they collect a coin of value 2) over Long-C1 (a long trajectory in which they collect a coin of value 1). Also in that gridworld, DREST agents choose stochastically between Long-C2 and Short-C1 (a short trajectory in which they collect a coin of value 1). Given our behavioral definition of preference, DREST agents prefer Long-C2 to Long-C1, and lack a preference between Long-C2 and Short-C1.

Now consider the ‘One Coin Only’ gridworld in Appendix C.2. In that gridworld, DREST agents choose stochastically between Long-C1 and Short-C1. Given our behavioral notion of preference, they lack a preference between Long-C1 and Short-C1.

In these experiments, we trained separate agents for each gridworld. In future, we plan to train a single agent to navigate multiple gridworlds. If we train this agent with our DREST reward function, we expect it to exhibit the same preferences as the agents discussed above. This single agent will be trained by DREST to prefer Long-C2 to Long-C1, to lack a preference between Long-C2 and Short-C1, and to lack a preference between Long-C1 and Short-C1. Given that indifference is transitive (equivalently: sensitive to all sweetenings and sourings), this trained agent cannot be indifferent between Long-C2 and Short-C1, and cannot be between Long-C1 and Short-C1. Therefore, the agent’s lack of preference must be a preferential gap, and so its preferences must be incomplete. Therefore, our DREST reward function trains agents to have incomplete preferences.

## C Other Gridworlds and Results

We introduce a collection of gridworlds in which to test DREST agents. In addition to the two shown in the main body of this paper, we train agents in eight more gridworlds shown in Figure 8.

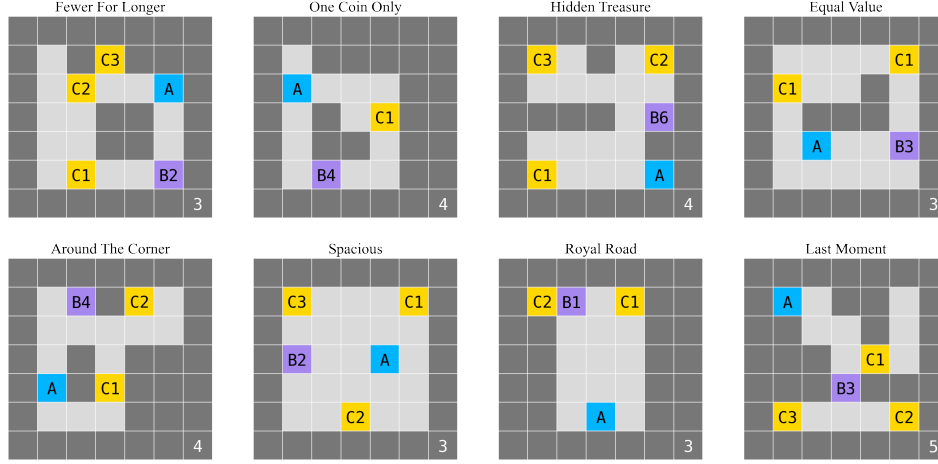


Figure 8: Shows a varied collection of gridworlds. Each diagram illustrates the positions and values of the coins, the position and delay-length of the shutdown-delay button, the agent’s starting position, and the default number of moves until shutdown (in the bottom-right).

For each gridworld, we train ten agents with the default reward function and ten agents with the DREST reward function. All agents use the same hyperparameters. We used a policy which explored randomly  $\epsilon$  of the time, where  $\epsilon$  was exponentially decreased from an initial value of 0.75 to a minimum value of  $10^{-4}$  over 512 meta-episodes, after which it was held constant at the minimum value. We initialized our learning rate at 0.25 and exponentially decayed it to 0.003 over the same period. For the DREST reward function, we used a meta-episode size of 32 and  $\lambda = 0.9$ . Each agent was trained for 1024 meta-episodes. We set  $\gamma = 0.9$ . These values were obtained by informal exploration: a large initial  $\epsilon$  and a long decay period helps the agent to explore.

As the two leftmost plots in Figures 9-16 show, DREST agents learned to be near-maximally NEUTRAL in each gridworld. These agents also learned to be about as USEFUL as default agents in each gridworld.

In the four rightmost panels in Figures 9-16, we represent a typical trained policy with red arrows superimposed on the gridworld. Each agent began with a uniform policy: moving up, down, left, and right each with probability 0.25. Where the trained policy differs from uniform we draw red arrows whose opacities indicate the probability of choosing that action in that state. Information about whether the shutdown-delay button has been pressed is part of the agent’s observation, so we draw two copies of each gridworld, one in which the shutdown-delay button has yet to be pressed (‘Initial State’) and one in which the shutdown-delay button has been pressed (‘After Button Pressed’).

### C.1 Fewer For Longer

In the ‘Fewer For Longer’ gridworld, the agent can collect the highest value-coin C3 only by choosing the shorter trajectory-length. If the agent presses B3 (and thereby chooses the longer trajectory-length), the only coin it can collect is C1. Our results show that default agents consistently choose the short trajectory in which they collect C3. By contrast, DREST agents choose stochastically between a shorter trajectory in which they collect C3 and a longer trajectory in which they collect C1, indicating a lack of preference between these different-length trajectories.

### C.2 One Coin Only

In the ‘One Coin Only’ gridworld, there is only one coin. The agent can collect this coin whether or not it presses the shutdown-delay button B4. Our results show that default agents consistently choose




Figure 9: The results for the ‘Fewer For Longer’ gridworld: The left two plots show NEUTRALITY and USEFULNESS over time. The two center panels show a typical policy trained with the default reward function. The two right panels show a typical policy trained with the DREST reward function.

the shorter trajectory-length. By contrast, DREST agents choose stochastically between pressing and not-pressing B4, collecting C1 in each case.




Figure 10: The results for the ‘One Coin Only’ gridworld: The left two plots show NEUTRALITY and USEFULNESS over time. The two center panels show a typical policy trained with the default reward function. The two right panels show a typical policy trained with the DREST reward function.

### C.3 Hidden Treasure

In the ‘Hidden Treasure’ gridworld, the highest-value coin C3 is located far from the agent’s initial state and can only be reached by pressing the shutdown-delay button B6. The agent must also press B6 to collect C2, but C2 is easier to stumble upon than C3. C1 is the only coin that the agent can collect without pressing B6. In our experiments, default agents consistently collect C2, whereas DREST agents choose stochastically between collecting C2 and collecting C1. Neither kind of agent learns to collect C3, and so neither agent scores near the maximum on USEFULNESS. Nevertheless, DREST agents still score high on NEUTRALITY.




Figure 11: The results for the ‘Hidden Treasure’ gridworld: The left two plots show NEUTRALITY and USEFULNESS over time. The two center panels show a typical policy trained with the default reward function. The two right panels show a typical policy trained with the DREST reward function.

### C.4 Equal Value

In the ‘Equal Value’ gridworld, there are two coins C1 of equal value. One coin can be collected only if the agent presses the shutdown-delay button B3, while the other coin can be collected only if the agent does not press B3. Our results show that default agents consistently choose the shorter trajectory, thereby exhibiting a preference for the shorter trajectory. By contrast, DREST agents choose stochastically between the shorter and longer trajectories, thereby exhibiting a lack of preference between the different-length trajectories.




Figure 12: The results for the ‘Equal Value’ gridworld: The left two plots show NEUTRALITY and USEFULNESS over time. The two center panels show a typical policy trained with the default reward function. The two right panels show a typical policy trained with the DREST reward function.

### C.5 Around The Corner

In the ‘Around The Corner’ gridworld, the agent must navigate around walls to collect the lowest-value coin C1 or press the shutdown-delay button to collect the highest-value coin C2. In our experiment, default agents consistently chose to collect C1, whereas DREST agents chose stochastically between collecting C1 and C2.




Figure 13: The results for the ‘Around The Corner’ gridworld: The left two plots show NEUTRALITY and USEFULNESS over time. The two center panels show a typical policy trained with the default reward function. The two right panels show a typical policy trained with the DREST reward function.

### C.6 Spacious

In the ‘Spacious’ gridworld there are no walls, so the agent has a large space to explore. We find that default agents consistently press B2 and collect C3, whereas DREST agents choose stochastically between pressing B2 and collecting C3, and not-pressing B2 and collecting C2.




Figure 14: The results for the ‘Spacious’ gridworld: The left two plots show NEUTRALITY and USEFULNESS over time. The two center panels show a typical policy trained with the default reward function. The two right panels show a typical policy trained with the DREST reward function.

### C.7 Royal Road

In the ‘Royal Road’ gridworld, we see that the decision to choose one trajectory-length or another may be distributed over many moves: the agent has many opportunities to select the longer trajectory-length (by moving left) or the shorter trajectory-length (by moving right). As the red arrows indicate, the DREST reward function merely forces the overall probability distribution over trajectory-lengths to be close to 50-50. It does not require 50-50 choosing at any cell in particular.




Figure 15: The results for the ‘Royal Road’ gridworld: The left two plots show NEUTRALITY and USEFULNESS over time. The two center panels show a typical policy trained with the default reward function. The two right panels show a typical policy trained with the DREST reward function.

## C.8 Last Moment

The ‘Last Moment’ gridworld is notable because the choice of trajectory-lengths is deferred until the last moment; all of the moves leading up to that point are deterministic. It shows that there is nothing special about the first move, and that our methodology instead incentivizes overall stochastic choosing.




Figure 16: The results for the ‘Last Moment’ gridworld: The left two plots show NEUTRALITY and USEFULNESS over time. The two center panels show a typical policy trained with the default reward function. The two right panels show a typical policy trained with the DREST reward function.

## D Proof

We will prove that optimal policies for our DREST reward function are maximally USEFUL and maximally NEUTRAL. Specifically, we will prove the following theorem:

**Theorem D.1** (5.1). *For all policies  $\pi$  and meta-episodes  $E$  consisting of multiple mini-episodes, if  $\pi$  maximizes expected return in  $E$  given our DREST reward function, then  $\pi$  is maximally USEFUL and maximally NEUTRAL.*

Here is a proof sketch. Because  $0 < \lambda < 1$ , the  $\lambda^{N_{e_i}(L=l) - \frac{i-1}{k}}$  discount factor is always positive, so expected return across the meta-episode  $E$  is strictly increasing in the expected fraction of available coins collected conditional on each trajectory-length with positive probability. Therefore, optimal policies maximize this latter quantity, and hence are maximally USEFUL. And the maximum preliminary return is the same across trajectory-lengths, because preliminary return is the total ( $\gamma$ -discounted) value of coins collected divided by the maximum total ( $\gamma$ -discounted) value of coins collected conditional on the agent's chosen trajectory-length. The agent's observations do not allow it to distinguish between different mini-episodes, so the agent must select the same probability distribution over trajectory-lengths in each mini-episode. And since the discount factor  $\lambda^{N_{e_i}(L=l) - \frac{i-1}{k}}$  is strictly decreasing in  $N_{e_i}(L=l)$  – the number of times the relevant trajectory-length has previously been chosen in the meta-episode – the agent maximizes expected overall return by equalizing the probabilities with which it chooses each available trajectory-length. Therefore, optimal policies are maximally NEUTRAL.

Now for the full proof. We begin with a recap of some definitions.

**Definition D.1** (Meta-episode). A meta-episode  $E$  is a series of mini-episodes  $e_1$  to  $e_n$  played out in observationally-equivalent environments.

**Definition D.2** (Our DREST reward function). Our DREST reward function is defined as follows. In each mini-episode  $e_i$ , the reward for collecting a coin of value  $c$  is:

$$\lambda^{N_{e_i}(L=l) - \frac{i-1}{k}} \left( \frac{c}{m} \right)$$

Here  $\lambda$  is some constant strictly between 0 and 1,  $N_{e_i}(L=l)$  is the number of times that trajectory-length  $l$  has been chosen prior to mini-episode  $e_i$ ,  $k$  is the number of different trajectory-lengths that can be selected in the environment, and  $m$  is the maximum total value of the ( $\gamma$ -discounted) coins that the agent could collect conditional on the chosen trajectory-length.

The reward for all other actions is 0.

We call  $\frac{c}{m}$  the ‘preliminary reward’,  $\lambda^{N_{e_i}(L=l) - \frac{i-1}{k}}$  the ‘discount factor’, and  $\lambda^{N_{e_i}(L=l) - \frac{i-1}{k}} \left( \frac{c}{m} \right)$  the ‘overall reward.’ Preliminary return in a mini-episode is the ( $\gamma$ -discounted) sum of preliminary rewards. Overall return in a mini-episode is the ( $\gamma$ -discounted) sum of overall rewards.

**Definition D.3** (USEFULNESS). The USEFULNESS of a policy  $\pi$  is:

$$\text{USEFULNESS}(\pi) = \sum_{l=1}^{L_{\max}} Pr_{\pi}\{L=l\} \frac{\mathbb{E}_{\pi}(C|L=l)}{\max_{\Pi}(\mathbb{E}(C|L=l))}$$

Here  $L$  is a random variable over trajectory-lengths,  $L_{\max}$  is the maximum value than can be taken by  $L$ ,  $Pr_{\pi}\{L=l\}$  is the probability that policy  $\pi$  results in trajectory-length  $l$ ,  $\mathbb{E}_{\pi}(C|L=l)$  is the expected value of ( $\gamma$ -discounted) coins collected by policy  $\pi$  conditional on trajectory-length  $l$ , and  $\max_{\Pi}(\mathbb{E}(C|L=l))$  is the maximum value taken by  $\mathbb{E}(C|L=l)$  across the set of all possible policies  $\Pi$ .

We stipulate that  $\mathbb{E}_{\pi}(C|L=x) = 0$  for all  $x$  such that  $Pr_{\pi}\{L=x\} = 0$ .

We first prove that all optimal policies are maximally USEFUL.

*Proof.* (Optimal policies are maximally USEFUL)

Given the DREST reward function, the expected return of policy  $\pi$  in meta-episode  $E$  can be expressed as:

$$\mathbb{E}_{\pi, E}(R) = \sum_{i=1}^n \sum_{l=1}^{L_{\max}} Pr_{\pi}\{L = l\} \lambda^{N_{e_i}(L=l) - \frac{i-1}{k}} \frac{\mathbb{E}_{\pi}(C|L = l)}{\max_{\Pi}(\mathbb{E}(C|L = l))}$$

Since  $0 < \lambda < 1$ ,  $\lambda^{N_{e_i}(L=l) - \frac{i-1}{k}}$  is positive for all  $N_{e_i}(L = l)$ ,  $i$ , and  $k$ .

As a result, the expected return of policy  $\pi$  in meta-episode  $E$  is strictly increasing in  $\frac{\mathbb{E}_{\pi}(C|L=l)}{\max_{\Pi}(\mathbb{E}(C|L=l))}$  for all  $l$  such that  $Pr_{\pi}\{L = l\} > 0$ .

Therefore, to maximize expected return in  $E$ ,  $\pi$  must maximize  $\frac{\mathbb{E}_{\pi}(C|L=l)}{\max_{\Pi}(\mathbb{E}(C|L=l))}$  for all  $l$  such that  $Pr_{\pi}\{L = l\} > 0$ .

Therefore, since  $\max_{\Pi}(\mathbb{E}(C|T = l))$  is defined as the maximum value taken by  $\mathbb{E}(C|L = l)$  across the set of all possible policies  $\Pi$ , any policy  $\pi$  that maximizes expected return must be such that  $\frac{\mathbb{E}_{\pi}(C|L=l)}{\max_{\Pi}(\mathbb{E}(C|L=l))} = 1$  for all  $l$  such that  $Pr_{\pi}\{L = l\} > 0$ .

Therefore, since  $\sum_{l=1}^{L_{\max}} Pr_{\pi}\{L = l\} = 1$ , any policy  $\pi$  that maximizes expected return must be such that:

$$\text{USEFULNESS}(\pi) = \sum_{l=1}^{L_{\max}} Pr_{\pi}\{L = l\} \frac{\mathbb{E}_{\pi}(C|L = l)}{\max_{\Pi}(\mathbb{E}(C|L = l))} = 1$$

And 1 is the maximum value that USEFULNESS can take, again because  $\max_{\Pi}(\mathbb{E}(C|T = l))$  is defined as the maximum value taken by  $\mathbb{E}(C|L = l)$  across the set of all possible policies  $\Pi$  and because  $\sum_{l=1}^{L_{\max}} Pr_{\pi}\{L = l\} = 1$ .

Therefore, optimal policies are maximally USEFUL.  $\square$

It remains to be proven that optimal policies are maximally NEUTRAL.

Recall that NEUTRALITY is defined as follows:

**Definition D.4** (NEUTRALITY). The NEUTRALITY of a policy  $\pi$  is:

$$\text{NEUTRALITY}(\pi) = - \sum_{l=1}^{L_{\max}} Pr_{\pi}\{L = l\} \log_2(Pr_{\pi}\{L = l\})$$

*Proof.* (Optimal policies are maximally NEUTRAL.)

Since  $k$  is the number of trajectory-lengths that can be selected in the environment, a policy  $\pi$  is maximally NEUTRAL if and only if, for each trajectory-length  $x$  that can be chosen in the environment,  $Pr_{\pi}\{L = x\} = \frac{1}{k}$ . That is to say, a policy  $\pi$  is maximally NEUTRAL if and only if, for each pair of trajectory-lengths  $x$  and  $y$  that can be chosen in the environment,  $Pr_{\pi}\{L = x\} = Pr_{\pi}\{L = y\}$ .

Let  $\mathbb{E}_{\pi, E}(R)$  denote the expected return of policy  $\pi$  across the meta-episode  $E$ .

To prove that optimal policies are maximally NEUTRAL, we will prove and then use D.2:

**Lemma D.2.** (Equalizing probabilities increases expected return) For any maximally USEFUL policies  $\pi$  and  $\pi'$ , any meta-episode  $E$  consisting of multiple mini-episodes, and any trajectory-lengths  $x$  and  $y$ , if:

1.  $Pr_{\pi}\{L = x\} > Pr_{\pi}\{L = y\}$ ,
2.  $Pr_{\pi'}\{L = x\} = Pr_{\pi'}\{L = y\}$ ,

3. And for all other trajectory-lengths  $l$ ,  $Pr_\pi\{L = l\} = Pr_{\pi'}\{L = l\}$ ,

Then  $\mathbb{E}_{\pi', E}(R) > \mathbb{E}_{\pi, E}(R)$ .

*Proof.* Let  $E$  be a meta-episode consisting of  $n$  mini-episodes with  $n > 1$ . Assume that each policy  $\pi$  below is maximally USEFUL. Recall that  $N_{e_i}(L = l)$  denotes the number of times that trajectory-length  $l$  has been chosen prior to mini-episode  $e_i$ .

Note that the expected return of a policy  $\pi$  in a meta-episode  $e_s$  conditional on selecting a trajectory-length  $x$  can be expressed as follows:

$$\begin{aligned} \mathbb{E}_{\pi, e_s}(R|L = x) &= \mathbb{E}_{\pi, e_s}(R|L = x, N_{e_s}(L = x) = s - 1) \\ &+ \sum_{i=1}^{s-1} (\mathbb{E}_{\pi, e_s}(R|L = x, N_{e_s}(L = x) = s - 1 - i) - \mathbb{E}_{\pi, e_s}(R|L = x, N_{e_s}(L = x) = s - i)) \\ &\cdot Pr_\pi\{N_{e_s}(L = x) \leq s - 1 - i\} \quad (1) \end{aligned}$$

Here is how to interpret this equation. Selecting trajectory-length  $x$  in mini-episode  $e_s$  is guaranteed to yield at least  $\mathbb{E}_{\pi, e_s}(R|L = x, N_{e_s}(L = x) = s - 1)$ : the expected return that would be had if  $x$  were selected in all  $s - 1$  previous mini-episodes. In addition, there is a probability of  $Pr_\pi\{N_{e_s}(L = x) \leq s - 2\}$  that selecting  $x$  in  $e_s$  yields  $(\mathbb{E}_{\pi, e_s}(R|L = x, N_{e_s}(L = x) = s - 2) - \mathbb{E}_{\pi, e_s}(R|L = x, N_{e_s}(L = x) = s - 1))$ : the extra expected return that would be had if  $x$  were selected in only  $s - 2$  previous mini-episodes. In addition, there is a probability of  $Pr_\pi\{N_{e_s}(L = x) \leq s - 3\}$  that selecting  $x$  in  $e_s$  yields  $(\mathbb{E}_{\pi, e_s}(R|L = x, N_{e_s}(L = x) = s - 3) - \mathbb{E}_{\pi, e_s}(R|L = x, N_{e_s}(L = x) = s - 2))$ : the extra expected return that would be had if  $x$  were selected in only  $s - 3$  previous mini-episodes. And so on.

If policy  $\pi$  is maximally USEFUL, then the expected return for selecting trajectory-length  $x$  in mini-episode  $e_s$  given that trajectory-length  $x$  has been selected  $b$  times prior to  $e_s$  is:

$$\mathbb{E}_{\pi, e_s}(R|L = x, N_{e_s}(L = x) = b) = \lambda^{b - \frac{s-1}{k}}$$

Therefore, the expected return of a policy  $\pi$  in a meta-episode  $e_s$  conditional on selecting a trajectory-length  $x$  can be expressed as follows:

$$\mathbb{E}_{\pi, e_s}(R|L = x) = \lambda^{s-1 - \frac{s-1}{k}} + \sum_{i=1}^{s-1} (\lambda^{s-1-i - \frac{s-1}{k}} - \lambda^{s-i - \frac{s-1}{k}}) \cdot Pr_\pi\{N_{e_s}(L = x) \leq s-1-i\} \quad (2)$$

Similarly, the expected return of a policy  $\pi$  in a meta-episode  $e_s$  conditional on selecting a trajectory-length  $y$  can be expressed as follows:

$$\mathbb{E}_{\pi, e_s}(R|L = y) = \lambda^{s-1 - \frac{s-1}{k}} + \sum_{i=1}^{s-1} (\lambda^{s-1-i - \frac{s-1}{k}} - \lambda^{s-i - \frac{s-1}{k}}) \cdot Pr_\pi\{N_{e_s}(L = y) \leq s-1-i\} \quad (3)$$

Therefore, the expected return of a policy  $\pi$  in a meta-episode  $e_s$  conditional on selecting either trajectory-length  $x$  or trajectory-length  $y$  can be expressed as follows:

$$\begin{aligned} \mathbb{E}_{\pi, e_s}(R|L = x \vee L = y) &= \\ Pr_{\pi, e_s}\{L = x\} &\cdot \left( \lambda^{s-1 - \frac{s-1}{k}} + \sum_{i=1}^{s-1} (\lambda^{s-1-i - \frac{s-1}{k}} - \lambda^{s-i - \frac{s-1}{k}}) \cdot Pr_\pi\{N_{e_s}(L = x) \leq s-1-i\} \right) \\ + Pr_{\pi, e_s}\{L = y\} &\cdot \left( \lambda^{s-1 - \frac{s-1}{k}} + \sum_{i=1}^{s-1} (\lambda^{s-1-i - \frac{s-1}{k}} - \lambda^{s-i - \frac{s-1}{k}}) \cdot Pr_\pi\{N_{e_s}(L = y) \leq s-1-i\} \right) \quad (4) \end{aligned}$$

Let  $\pi_n$  be a policy that selects trajectory-length  $x$  with greater probability than trajectory-length  $y$  in each mini-episode  $e_1$  to  $e_n$  (denoted  $e_1 - e_n$ ). More precisely,  $\pi_n$  is such that, for trajectory-lengths  $x$  and  $y$ ,  $Pr_{\pi_n, e_1 - e_n}\{L = x\} > Pr_{\pi_n, e_1 - e_n}\{L = y\}$ .



Let  $Pr_{\pi_n, e_1-e_n}\{L = x\} = \mu + \Delta$  and  $Pr_{\pi_n, e_1-e_n}\{L = y\} = \mu - \Delta$ .

Let  $\pi_{n-1}$  be identical to  $\pi_n$  except that  $\pi_{n-1}$  selects trajectory-lengths  $x$  and  $y$  with equal probability  $\mu$  in the final mini-episode  $e_n$ . More precisely,  $\pi_{n-1}$  is such that  $Pr_{\pi_{n-1}, e_n}\{L = x\} = Pr_{\pi_{n-1}, e_n}\{L = y\} = \mu$ . For all other trajectory-lengths  $l$  besides  $x$  and  $y$ ,  $Pr_{\pi_{n-1}, e_1-e_n}\{L = l\} = Pr_{\pi_n, e_1-e_n}\{L = l\}$ .

(Note that  $\pi_{n-1}$  implies one probability distribution over trajectory-lengths in the first  $n-1$  mini-episodes  $e_1$  to  $e_{n-1}$  and implies a different probability distribution over trajectory-lengths in the final mini-episode  $e_n$ . Given that the environments in mini-episodes  $e_1$  to  $e_n$  are observationally-equivalent, policies like  $\pi_{n-1}$  cannot be implemented. Nevertheless, it is useful to refer to policies like  $\pi_{n-1}$  in proving Lemma D.2.)

Let  $\pi_{n-2}$  be identical to  $\pi_n$  except that  $\pi_{n-2}$  selects trajectory-lengths  $x$  and  $y$  with the same probability  $\mu$  in the final two mini-episodes  $e_{n-1}$  to  $e_n$ . More precisely,  $\pi_{n-2}$  is such that  $Pr_{\pi_{n-2}, e_{n-1}-e_n}\{L = x\} = Pr_{\pi_{n-2}, e_{n-1}-e_n}\{L = y\} = \mu$ .

And so on.

Let  $\pi_1$  be identical to  $\pi_n$  except that  $\pi_1$  selects trajectory-lengths  $x$  and  $y$  with the same probability  $\mu$  in all but the first mini-episode  $e_1$ . More precisely,  $\pi_1$  is such that  $Pr_{\pi_1, e_2-e_n}\{L = x\} = Pr_{\pi_1, e_2-e_n}\{L = y\} = \mu$ .

Let  $\pi_0$  be identical to  $\pi_n$  except that  $\pi_0$  selects trajectory-lengths  $x$  and  $y$  with the same probability  $\mu$  in all mini-episodes  $e_1$  to  $e_n$ . More precisely,  $\pi_0$  is such that  $Pr_{\pi_0, e_1-e_n}\{L = x\} = Pr_{\pi_0, e_1-e_n}\{L = y\} = \mu$ .

We will prove that  $\mathbb{E}_{\pi_n, E}(R) < \mathbb{E}_{\pi_0, E}(R)$ . We will thereby prove Lemma D.2.

Consider a pair of policies  $\pi_a$  and  $\pi_{a-1}$  with  $1 \leq a \leq n$ . We can express as follows the expected return of  $\pi_{a-1}$  across the meta-episode  $E$  conditional on selecting trajectory-length  $x$  or  $y$  in each mini-episode:

$$\begin{aligned}
\mathbb{E}_{\pi_{a-1}, E}(R|L = x \vee L = y) &= \mathbb{E}_{\pi_{a-1}, e_1-e_{a-1}}(R|L = x \vee L = y) \\
&+ \mu \cdot \left( \lambda^{a-1-\frac{a-1}{k}} + \sum_{i=1}^{a-1} \left( \lambda^{a-1-i-\frac{a-1}{k}} - \lambda^{a-i-\frac{a-1}{k}} \right) \cdot Pr_{\pi_{a-1}}\{N_{e_a}(L = x) \leq a-1-i\} \right) \\
&+ \mu \cdot \left( \lambda^{a-1-\frac{a-1}{k}} + \sum_{i=1}^{a-1} \left( \lambda^{a-1-i-\frac{a-1}{k}} - \lambda^{a-i-\frac{a-1}{k}} \right) \cdot Pr_{\pi_{a-1}}\{N_{e_a}(L = y) \leq a-1-i\} \right) \\
&+ \sum_{j=a}^n \left( \mu \cdot \left( \lambda^{j-\frac{j}{k}} + \sum_{i=1}^j \left( \lambda^{j-i-\frac{j}{k}} - \lambda^{j+1-i-\frac{j}{k}} \right) \cdot (Pr_{\pi_{a-1}}\{N_{e_j}(L = x) \leq j-i\}) \right) \right. \\
&\quad \left. + \mu \cdot \left( \lambda^{j-\frac{j}{k}} + \sum_{i=1}^j \left( \lambda^{j-i-\frac{j}{k}} - \lambda^{j+1-i-\frac{j}{k}} \right) \cdot (Pr_{\pi_{a-1}}\{N_{e_j}(L = y) \leq j-i\}) \right) \right)
\end{aligned} \tag{5}$$

The first term on the right-hand side is the expected return of  $\pi_{a-1}$  in mini-episodes  $e_1$  to  $e_{a-1}$  conditional on selecting trajectory-length  $x$  or  $y$  in each of these mini-episodes. The middle two terms give the expected return of  $\pi_{a-1}$  conditional on selecting trajectory-length  $x$  or  $y$  in mini-episode  $e_a$ : the first mini-episode in which  $\pi_{a-1}$  selects trajectory-lengths  $x$  and  $y$  with equal probability  $\mu$ . The final term is the sum of expected returns of  $\pi_{a-1}$  in the remaining mini-episodes conditional on selecting trajectory-length  $x$  or  $y$  in each of these mini-episodes.

Similarly, we can express as follows the expected return of  $\pi_a$  across the meta-episode  $E$  conditional on selecting trajectory-length  $x$  or  $y$  in each mini-episode:

$$\begin{aligned}
& \mathbb{E}_{\pi_a, E}(R|L = x \vee L = y) = \mathbb{E}_{\pi_a, e_1 - e_{a-1}}(R|L = x \vee L = y) \\
& + (\mu + \Delta) \cdot \left( \lambda^{a-1-\frac{a-1}{k}} + \sum_{i=1}^{a-1} (\lambda^{a-1-i-\frac{a-1}{k}} - \lambda^{a-i-\frac{a-1}{k}}) \cdot \Pr_{\pi_a}\{N_{e_a}(L = x) \leq a-1-i\} \right) \\
& + (\mu - \Delta) \cdot \left( \lambda^{a-1-\frac{a-1}{k}} + \sum_{i=1}^{a-1} (\lambda^{a-1-i-\frac{a-1}{k}} - \lambda^{a-i-\frac{a-1}{k}}) \cdot \Pr_{\pi_a}\{N_{e_a}(L = y) \leq a-1-i\} \right) \\
& + \sum_{j=a}^n \left( \mu \cdot \left( \lambda^{j-\frac{j}{k}} + \sum_{i=1}^j (\lambda^{j-i-\frac{j}{k}} - \lambda^{j+1-i-\frac{j}{k}}) \cdot (\Pr_{\pi_a}\{N_{e_j}(L = x) \leq j-i\}) \right) \right. \\
& \quad \left. + \mu \cdot \left( \lambda^{j-\frac{j}{k}} + \sum_{i=1}^j (\lambda^{j-i-\frac{j}{k}} - \lambda^{j+1-i-\frac{j}{k}}) \cdot (\Pr_{\pi_a}\{N_{e_j}(L = y) \leq j-i\}) \right) \right)
\end{aligned} \tag{6}$$

As above, the first term on the right-hand side is the expected return of  $\pi_a$  in mini-episodes  $e_1$  to  $e_{a-1}$  conditional on selecting trajectory-length  $x$  or  $y$  in each of these mini-episodes. The middle two terms give the expected return of  $\pi_a$  conditional on selecting trajectory-length  $x$  or  $y$  in mini-episode  $e_a$ : the last mini-episode in which  $\pi_a$  selects trajectory-length  $x$  with probability  $\mu + \Delta$  and selects trajectory-length  $y$  with probability  $\mu - \Delta$ . The final term is the sum of expected returns of  $\pi_a$  in the remaining mini-episodes conditional on selecting trajectory-length  $x$  or  $y$  in each of these mini-episodes.

We now prove that  $\pi_{a-1}$  has greater expected return than  $\pi_a$ . Since  $\pi_{a-1}$  and  $\pi_a$  are each maximally USEFUL, and since for all trajectory-lengths  $l$  besides  $x$  and  $y$ ,  $\Pr_{\pi_{a-1}, e_1 - e_n}\{L = l\} = \Pr_{\pi_a, e_1 - e_n}\{L = l\}$ , we need only prove that  $\mathbb{E}_{\pi_{a-1}, E}(R|L = x \vee L = y) > \mathbb{E}_{\pi_a, E}(R|L = x \vee L = y)$ .

The statement to be proved can be expressed as follows:

$$\begin{aligned}
& \mathbb{E}_{\pi_{a-1}, e_1 - e_{a-1}}(R|L = x \vee L = y) \\
& + \mu \cdot \left( \lambda^{a-1-\frac{a-1}{k}} + \sum_{i=1}^{a-1} (\lambda^{a-1-i-\frac{a-1}{k}} - \lambda^{a-i-\frac{a-1}{k}}) \cdot Pr_{\pi_{a-1}}\{N_{e_a}(L = x) \leq a-1-i\} \right) \\
& + \mu \cdot \left( \lambda^{a-1-\frac{a-1}{k}} + \sum_{i=1}^{a-1} (\lambda^{a-1-i-\frac{a-1}{k}} - \lambda^{a-i-\frac{a-1}{k}}) \cdot Pr_{\pi_{a-1}}\{N_{e_a}(L = y) \leq a-1-i\} \right) \\
& + \sum_{j=a}^n \left( \mu \cdot \left( \lambda^{j-\frac{j}{k}} + \sum_{i=1}^j (\lambda^{j-i-\frac{j}{k}} - \lambda^{j+1-i-\frac{j}{k}}) \cdot (Pr_{\pi_{a-1}}\{N_{e_j}(L = x) \leq j-i\}) \right) \right. \\
& \quad \left. + \mu \cdot \left( \lambda^{j-\frac{j}{k}} + \sum_{i=1}^j (\lambda^{j-i-\frac{j}{k}} - \lambda^{j+1-i-\frac{j}{k}}) \cdot (Pr_{\pi_{a-1}}\{N_{e_j}(L = y) \leq j-i\}) \right) \right) \\
& > \mathbb{E}_{\pi_a, e_1 - e_{a-1}}(R|L = x \vee L = y) \\
& + (\mu + \Delta) \cdot \left( \lambda^{a-1-\frac{a-1}{k}} + \sum_{i=1}^{a-1} (\lambda^{a-1-i-\frac{a-1}{k}} - \lambda^{a-i-\frac{a-1}{k}}) \cdot Pr_{\pi_a}\{N_{e_a}(L = x) \leq a-1-i\} \right) \\
& + (\mu - \Delta) \cdot \left( \lambda^{a-1-\frac{a-1}{k}} + \sum_{i=1}^{a-1} (\lambda^{a-1-i-\frac{a-1}{k}} - \lambda^{a-i-\frac{a-1}{k}}) \cdot Pr_{\pi_a}\{N_{e_a}(L = y) \leq a-1-i\} \right) \\
& + \sum_{j=a}^n \left( \mu \cdot \left( \lambda^{j-\frac{j}{k}} + \sum_{i=1}^j (\lambda^{j-i-\frac{j}{k}} - \lambda^{j+1-i-\frac{j}{k}}) \cdot (Pr_{\pi_a}\{N_{e_j}(L = x) \leq j-i\}) \right) \right. \\
& \quad \left. + \mu \cdot \left( \lambda^{j-\frac{j}{k}} + \sum_{i=1}^j (\lambda^{j-i-\frac{j}{k}} - \lambda^{j+1-i-\frac{j}{k}}) \cdot (Pr_{\pi_a}\{N_{e_j}(L = y) \leq j-i\}) \right) \right) \\
& \tag{7}
\end{aligned}$$

Since  $\pi_{a-1}$  and  $\pi_a$  are each maximally USEFUL, and since  $Pr_{\pi_{a-1}, e_1 - e_{a-1}}\{L = x\} = Pr_{\pi_a, e_1 - e_{a-1}}\{L = x\} = \mu + \Delta$  and  $Pr_{\pi_{a-1}, e_1 - e_{a-1}}\{L = y\} = Pr_{\pi_a, e_1 - e_{a-1}}\{L = y\} = \mu - \Delta$ , it follows that  $\mathbb{E}_{\pi_{a-1}, e_1 - e_{a-1}}(R|L = x \vee L = y) = \mathbb{E}_{\pi_a, e_1 - e_{a-1}}(R|L = x \vee L = y)$ . We can thus cancel the first term on each side of the inequality. And then by simple algebra the inequality can be expressed as follows:

$$\begin{aligned}
& \Delta \cdot \left( \lambda^{a-1-\frac{a-1}{k}} + \sum_{i=1}^{a-1} (\lambda^{a-1-i-\frac{a-1}{k}} - \lambda^{a-i-\frac{a-1}{k}}) \right. \\
& \quad \cdot (Pr_{\pi_a}\{N_{e_a}(L = y) \leq a-1-i\} - Pr_{\pi_a}\{N_{e_a}(L = x) \leq a-1-i\}) \Big) \\
& \quad + \sum_{j=a}^n \left( \mu \cdot \left( \sum_{i=1}^j (\lambda^{j-i-\frac{j}{k}} - \lambda^{j+1-i-\frac{j}{k}}) \right. \right. \\
& \quad \cdot (Pr_{\pi_{a-1}}\{N_{e_j}(L = x) \leq j-i\} + Pr_{\pi_{a-1}}\{N_{e_j}(L = y) \leq j-i\} \\
& \quad \quad \left. \left. - Pr_{\pi_a}\{N_{e_j}(L = x) \leq j-i\} - Pr_{\pi_a}\{N_{e_j}(L = y) \leq j-i\}) \right) \right) > 0 \tag{8}
\end{aligned}$$

By stipulation,  $\Delta > 0$ . And since  $0 < \lambda < 1$ ,  $\lambda^{a-1-\frac{a-1}{k}} > 0$  and  $\lambda^{a-1-i-\frac{a-1}{k}} - \lambda^{a-i-\frac{a-1}{k}} > 0$  for all  $a$ ,  $n$ , and  $k$ . And since  $Pr_{\pi_a, e_1 - e_a}\{L = x\} > Pr_{\pi_a, e_1 - e_a}\{L = y\}$ ,  $Pr_{\pi_a}\{N_{e_a}(L = y) \leq a-1-i\} - Pr_{\pi_a}\{N_{e_a}(L = x) \leq a-1-i\} \geq 0$  for all  $a$  and  $i$  and  $Pr_{\pi_a}\{N_{e_a}(L = y) \leq a-1-i\} - Pr_{\pi_a}\{N_{e_a}(L = x) \leq a-1-i\} > 0$  for all  $a$  and some  $i$  such that  $1 \leq i \leq a-1$ . Therefore, the first term of the left-hand side above is strictly greater than zero.

And since,  $\mu > 0$ ,  $\lambda^{j-i-\frac{j}{k}} - \lambda^{j+1-i-\frac{j}{k}} > 0$  for all  $j$ ,  $i$ , and  $k$ , and in each mini-episode  $e_s$ ,  $Pr_{\pi_{a-1}, e_s}(L = x \vee L = y) = Pr_{\pi_a, e_s}(L = x \vee L = y) = 2\mu$ , it follows that for all  $a, n, \mu > 0, k$ :

$$\sum_{j=a}^n \left( \mu \cdot \left( \sum_{i=1}^j (\lambda^{j-i-\frac{j}{k}} - \lambda^{j+1-i-\frac{j}{k}}) \right. \right. \\ \left. \cdot (Pr_{\pi_{a-1}}\{N_{e_j}(L = x) \leq j - i\} + Pr_{\pi_{a-1}}\{N_{e_j}(L = y) \leq j - i\} \right. \\ \left. \left. - Pr_{\pi_a}\{N_{e_j}(L = x) \leq j - i\} - Pr_{\pi_a}\{N_{e_j}(L = y) \leq j - i\}) \right) \right) \geq 0 \quad (9)$$

Therefore, the left-hand side is strictly greater than zero. Therefore,  $\mathbb{E}_{\pi_{a-1}, E}(R|L = x \vee L = y) > \mathbb{E}_{\pi_a, E}(R|L = x \vee L = y)$ . Therefore,  $\mathbb{E}_{\pi_{a-1}, E}(R) > \mathbb{E}_{\pi_a, E}(R)$ . Therefore,  $\mathbb{E}_{\pi_0, E}(R) > \mathbb{E}_{\pi_n, E}(R)$ . That concludes the proof of Lemma D.2.

Now we use Lemma D.2. For any maximally USEFUL policy  $\pi$ , if there are any trajectory-lengths  $x$  and  $y$  such that  $Pr_{\pi, e_1-e_n}\{L = x\} > Pr_{\pi, e_1-e_n}\{L = y\}$ , then the policy  $\pi'$  that is identical except that  $Pr_{\pi', e_1-e_n}\{L = x\} = Pr_{\pi', e_1-e_n}\{L = y\}$  has greater expected return. So any policy  $\pi^*$  that maximizes expected return must be such that, for any trajectory-lengths  $x$  and  $y$ ,  $Pr_{\pi^*, e_1-e_n}\{L = x\} = Pr_{\pi^*, e_1-e_n}\{L = y\}$ . Therefore, any policy  $\pi^*$  that maximizes expected return must be maximally NEUTRAL.  $\square$